

CS 235: Logical Agent


Propositional Logic

Dr. Moumita Roy
CSE Dept., IIITG

Agents

- Definition: An **agent** perceives its **environment** via **sensors** and acts upon that environment through its

a



Examples of Agent

- **A software agent** has Keystrokes, file contents, received network packages which act as sensors and displays on the screen, files, sent network packets acting as actuators
- **A Human agent** has eyes, ears, and other organs which act as sensors and hands, legs, mouth, and other body parts acting as actuators
- **A Robotic agent** has Cameras and infrared range finders which act as sensors and various motors acting as actuators

Rational agents

- An agent should strive to "do the right thing", based on what:
 - it can perceive and
 - the actions it can perform.
- **Right action:** Select the action which can maximize the performance

Performance measure: *An objective criterion for success of an agent's behavior.*

Performance measures of a vacuum-cleaner agent: amount of dirt cleaned up, amount of time taken, amount of electricity consumed, level of noise generated, etc.

Performance measures self-driving car: time to reach destination (minimize), safety, predictability of behavior for other agents, reliability, etc.

Performance measure of game-playing agent: win/loss percentage (maximize), robustness, unpredictability (to "confuse" opponent), etc.

Types of AI agents


Table-lookup driven agents

- Uses a percept sequence / action table in memory to
- find the next action. Implemented as a (large) lookup table.

Drawbacks:

- Huge table (often simply too large)
- Takes a long time to build/learn the table

Toy example:
Vacuum world.



Percepts: robot senses it's **location** and “**cleanliness.**”

So, **location** and **contents**, e.g., [A, Dirty], [B, Clean].

With 2 locations, we get **4 different possible sensor inputs.**

Actions: *Left, Right, Suck, NoOp*

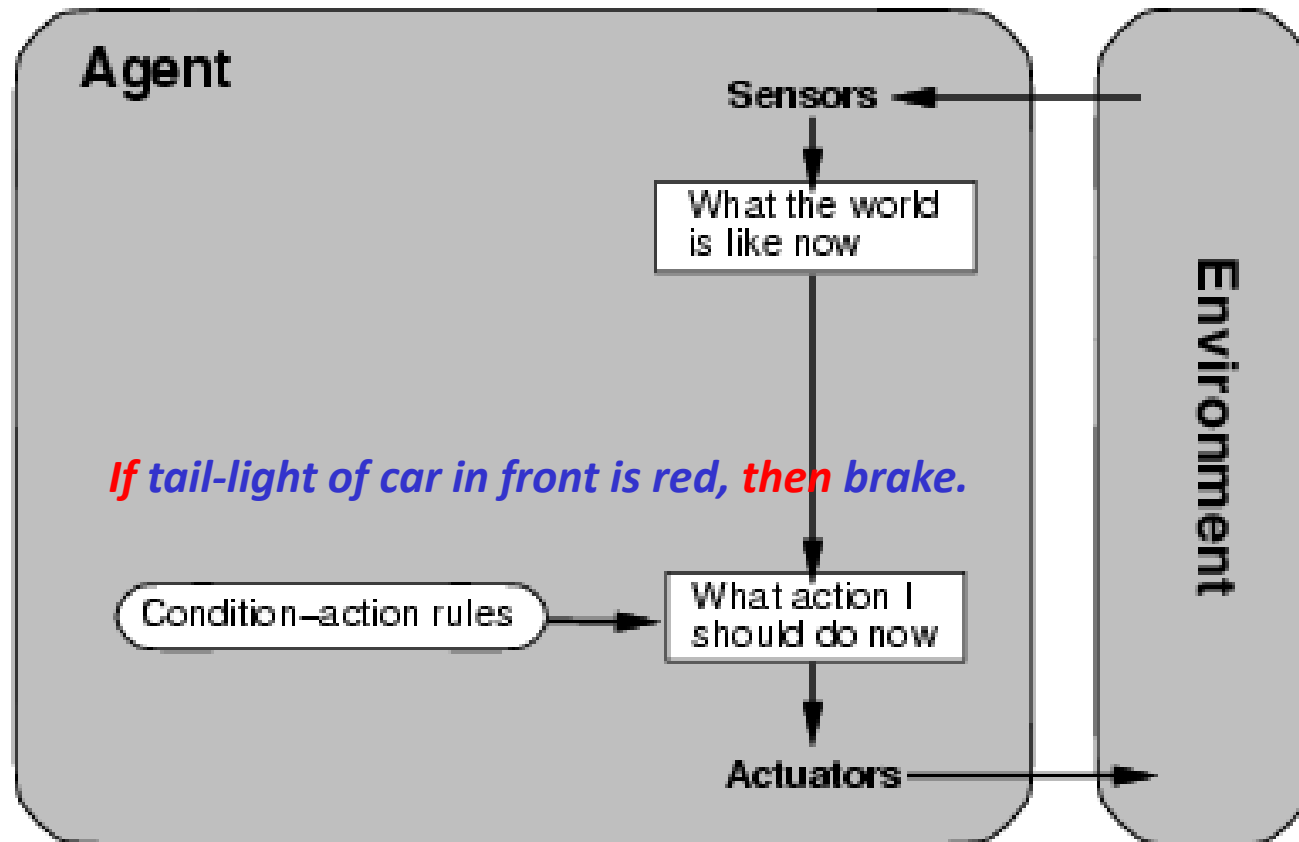
Percept sequence	Action
[A, Clean]	<i>Right</i>
[A, Dirty]	<i>Suck</i>
[B, Clean]	<i>Left</i>
[B, Dirty]	<i>Suck</i>
[A, Clean], [A, Clean]	<i>Right</i>
[A, Clean], [A, Dirty]	<i>Suck</i>
⋮	⋮
[A, Clean], [A, Clean], [A, Clean]	<i>Right</i>
[A, Clean], [A, Clean], [A, Dirty]	<i>Suck</i>
⋮	⋮

Figure 2.3 Partial tabulation of a simple agent function for the vacuum-cleaner world shown in Figure 2.2.

Simple reflex agents

- The Simple reflex agents are the simplest agents. These agents take decisions on the basis of the current percepts and ignore the rest of the percept history
- These agents only succeed in the fully observable environment
- The Simple reflex agent does not consider any part of percepts history during their decision and action process
- The Simple reflex agent works on **Condition-action rule**, which means it maps the current state to action
- Problems for the simple reflex agent design approach:
 - They have very limited intelligence
 - They do not have knowledge of non-perceptual parts of the current state
 - Mostly too big to generate and to store.
 - Not adaptive to changes in the environment.

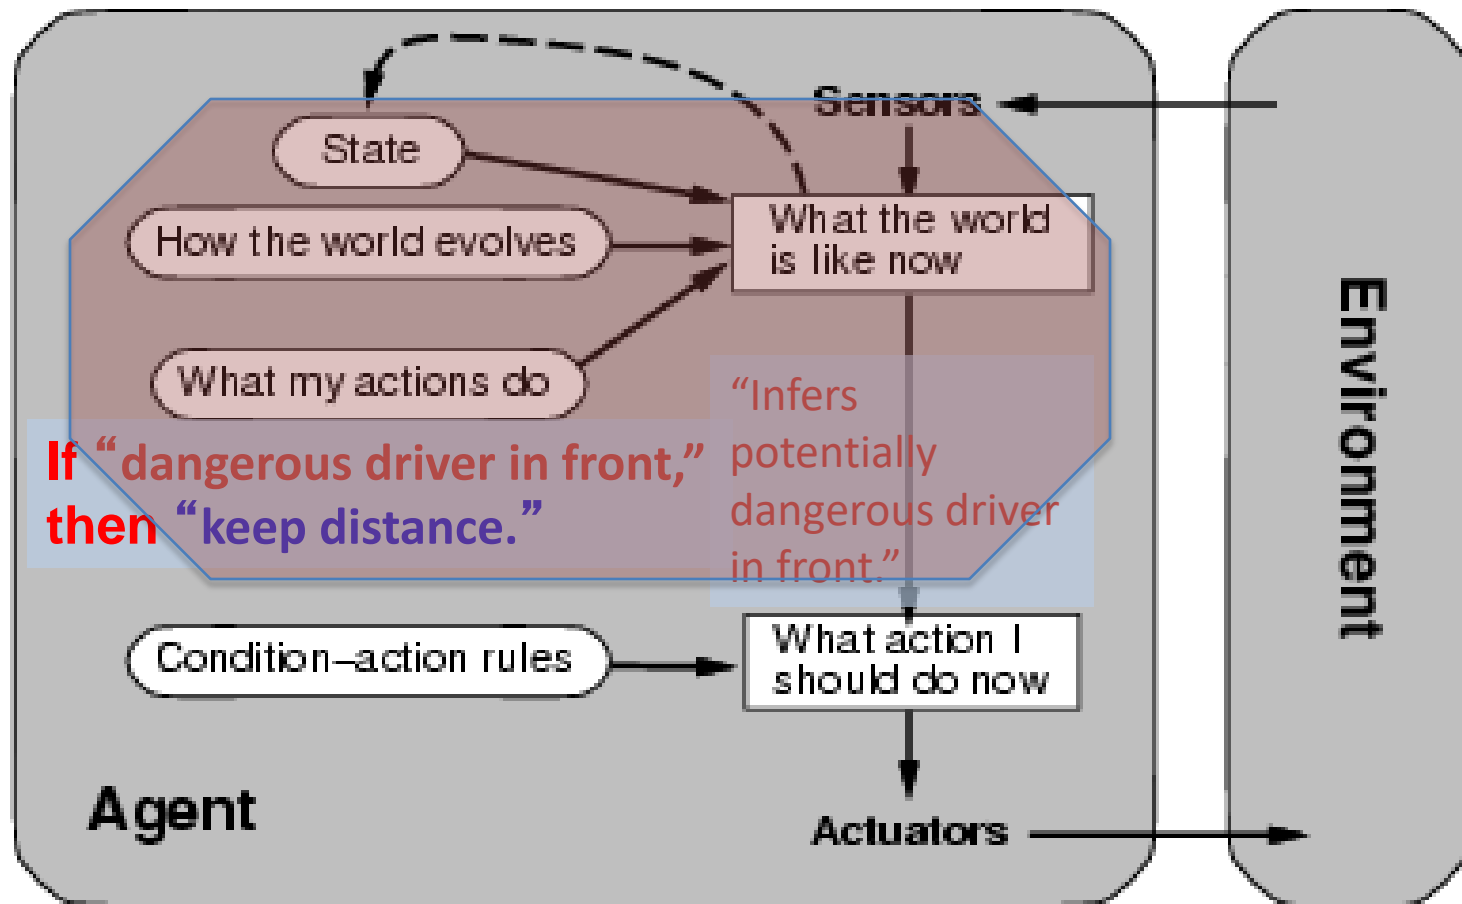
Agent selects actions on the basis
of *current percept only*.



Model-based reflex agents

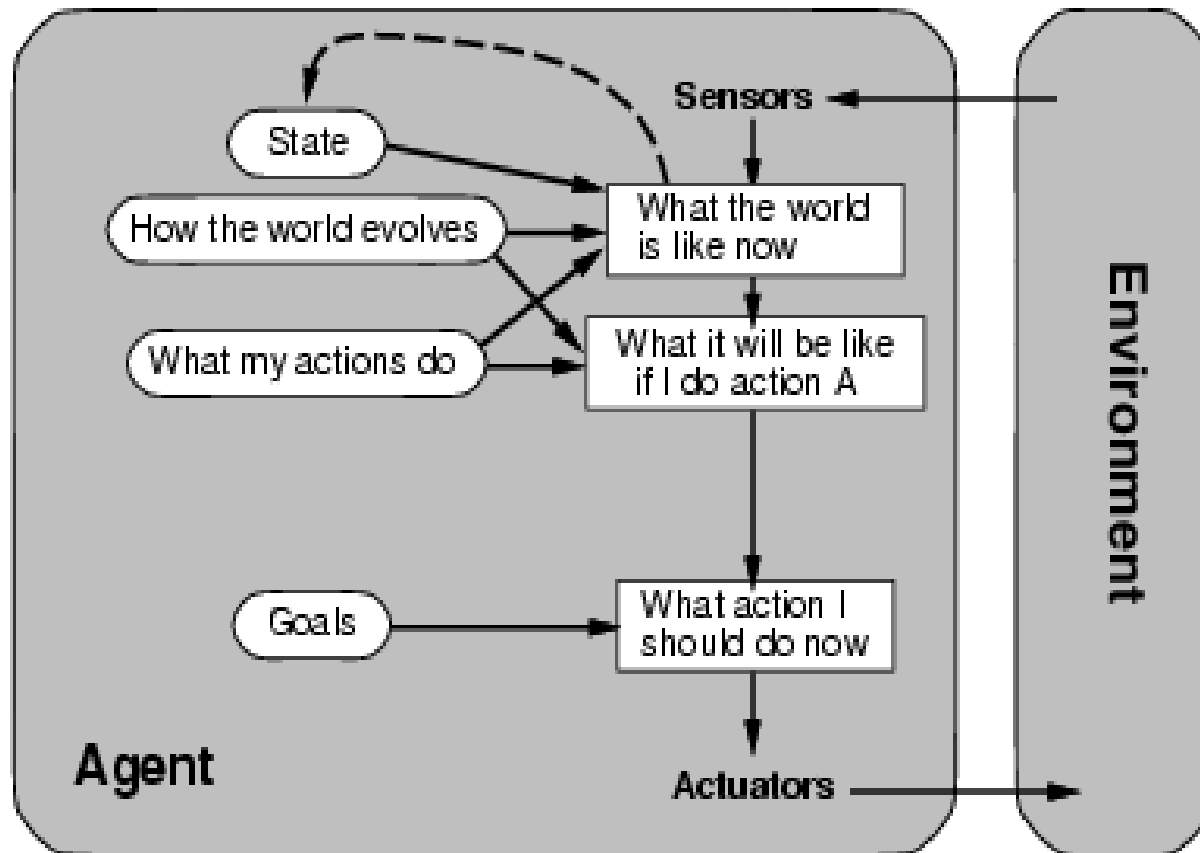
- The Model-based agent can work in a partially observable environment, and track the situation
- A model-based agent has two important factors:
 - **Model:** It is knowledge about "how things happen in the world," so it is called a Model-based agent
 - **Internal State:** It is a representation of the current state based on percept history
- These agents have the model, "which is knowledge of the world" and based on the model they perform actions
- Updating the agent state requires information about:
 - How the world evolves
 - How the agent's action affects the world

Model-based reflex agents



Goal-based agents

- The knowledge of the current state environment is not always sufficient to decide for an agent to what to do
 - The agent needs to know its goal which describes desirable situations
 - Goal-based agents expand the capabilities of the model-based agent by having the "goal" information
 - They choose an action, so that they can achieve the goal
 - These agents may have to consider a long sequence of possible actions before deciding whether the goal is achieved or not
 - Such considerations of different scenario are called **searching and planning**, which makes an agent proactive
- **problem solving and search!**



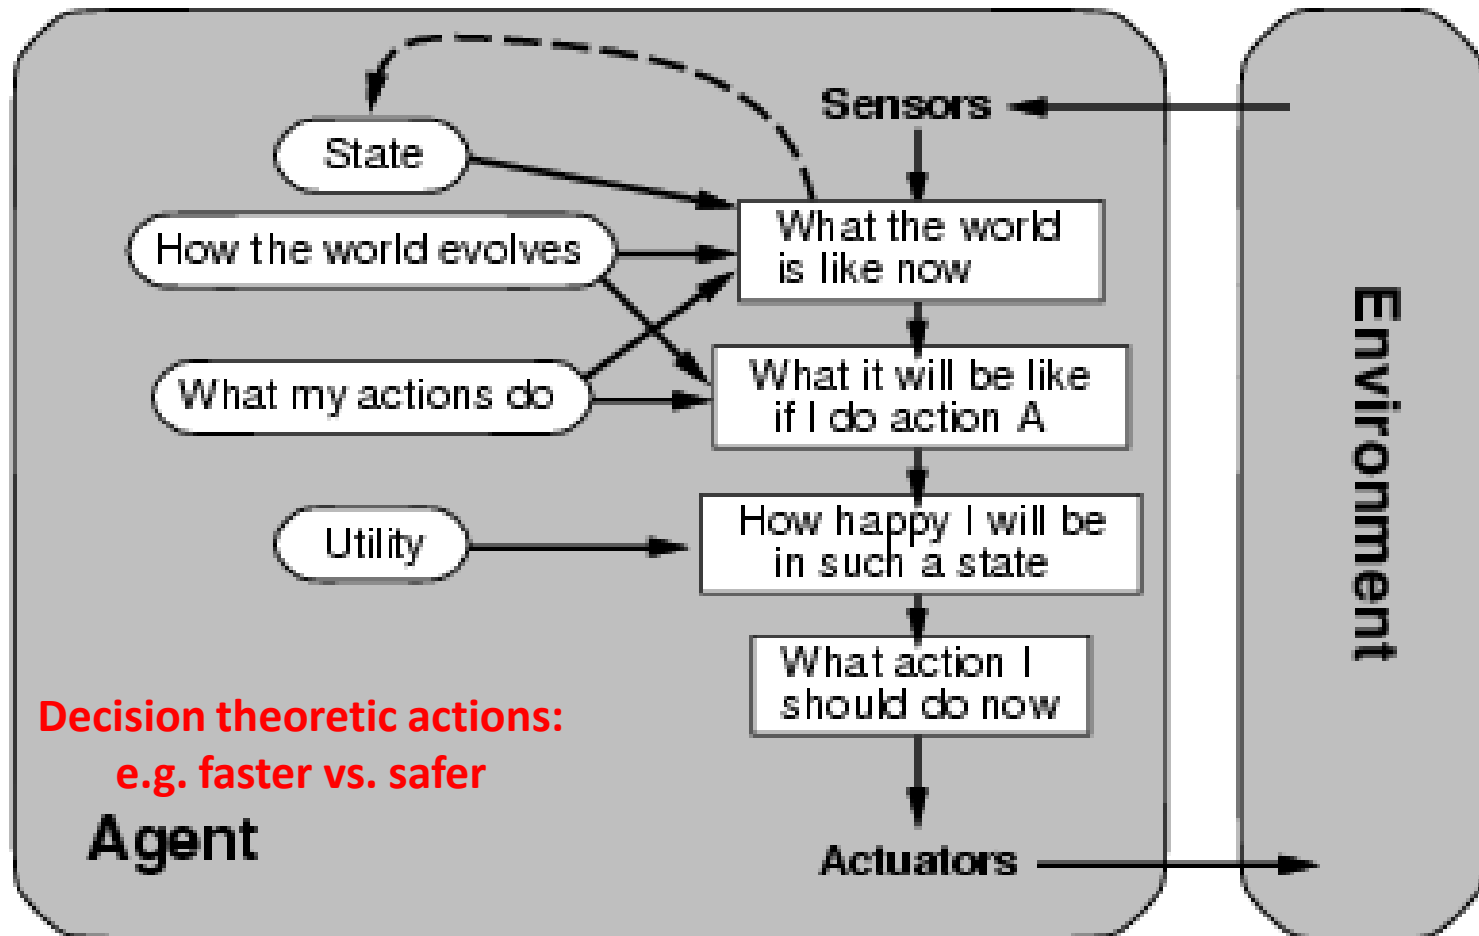
Agent keeps track of the world state as well as set of goals it's trying to achieve: chooses actions that will (eventually) lead to the goal(s).

More flexible than reflex agents → may involve search and planning

Utility-based agents

- These agents are similar to the goal-based agent but provide an extra component of utility measurement which makes them different by providing a measure of success at a given state
- Utility-based agent act based not only goals but also the best way to achieve the goal
- The Utility-based agent is useful when there are multiple possible alternatives, and an agent has to choose in order to perform the best action
- The utility function maps each state to a real number to check how efficiently each action achieves the goals

Utility-based agents




Learning Agent

- A learning agent in AI is the type of agent which can learn from its past experiences, or it has learning capabilities
- It starts to act with basic knowledge and then able to act and adapt automatically through learning.
- Hence, learning agents are able to learn, analyze performance, and look for new ways to improve the performance.


Model-based reflex agents/Logical Agent

- Human do things not only purely on reflex mechanism; but by the process of reasoning that operate on the internal representation of the knowledge
- To address these issues we will introduce:
 - A knowledge base (KB): a list of facts that are known to the agent; set of sentences in a formal language
 - Rules to infer new facts from old facts using rules of inference
 - Logic provides the formal language for this
- Declarative approach to building an agent:
 - Tell it what it needs to know.
 - Ask it what to do -> answers should follow from the KB.

Agent Architecture: Logical Agents



(a) Atomic



(b) Factored

A model is a **structured** representation of the world

- Graph-Based Search: State is **black box**, no internal structure, atomic
- Factored Representation: State is list or vector of facts
- Facts are expressed in **formal logic**.

Knowledge-Based Agents

- KB = knowledge base
 - A set of sentences or facts
 - e.g., a set of statements in a logic language
- Inference
 - Deriving new sentences from old
 - e.g., using a set of logical statements to infer new ones
- A simple model for reasoning
 - Agent is told or perceives new evidence
 - E.g., A is true
 - Agent then infers new facts to add to the KB
 - E.g., $KB = \{ A \rightarrow (B \text{ OR } C) \}$, then given A and not C we can infer that B is true
 - B is now added to the KB even though it was not explicitly asserted, i.e., the agent inferred B

Logic

- **Logics** are formal languages for representing information such that conclusions can be drawn
- **Syntax** defines the sentences in the language
- **Semantics** define the "meaning" of sentences;
 - i.e., define **truth** of a sentence in a world
- E.g., the language of arithmetic
 - $x+2 \geq y$ is a sentence; $x^2+y > \{\}$ is not a sentence \longrightarrow syntax
 -
 - $x+2 \geq y$ is true in a world where $x = 7, y = 1$
 - $x+2 \geq y$ is false in a world where $x = 0, y = 6$ } semantics

Entailment

- **Entailment** means that one thing **follows from** another:

$$KB \models \alpha$$

- Knowledge base KB entails sentence α if and only if α is true in **all worlds** where KB is true
 - E.g., the KB containing “the students of B.Tech won and the students of M.Tech won” entails “The students of B.Tech won”.

To sum up

- A formal language
 - KB = set of sentences
- Syntax
 - what sentences are legal (well-formed)
 - E.g., arithmetic
 - $X+2 \geq y$ is a wf sentence, $+x2y$ is not a wf sentence
- Semantics
 - loose meaning: the interpretation of each sentence
 - More precisely:
 - Defines the truth of each sentence wrt to each possible world
 - e.g.,
 - $X+2 = y$ is true in a world where $x=7$ and $y=9$
 - $X+2 = y$ is false in a world where $x=7$ and $y=1$
 - Note: standard logic - each sentence is T or F wrt each world
 - Fuzzy logic - allows for degrees of truth.

Propositional logic: Syntax

- Propositional logic is the simplest logic - illustrates basic ideas
- Atomic sentences = single proposition symbols
 - E.g., P , Q , R (begin with capital letter)
 - Two distinguished atom: True , False
- Complex sentences:
 - If S is a sentence, $\neg S$ is a sentence (negation)
 - If S_1 and S_2 are sentences, $S_1 \wedge S_2$ is a sentence (conjunction)
 - If S_1 and S_2 are sentences, $S_1 \vee S_2$ is a sentence (disjunction)
 - If S_1 and S_2 are sentences, $S_1 \Rightarrow S_2$ is a sentence (implication)
 - If S_1 and S_2 are sentences, $S_1 \Leftrightarrow S_2$ is a sentence (biconditional)

Propositional logic

Consider the propositional logic where

P: It is hot

Q: It is humid

R: It is raining

Formalized the following languages:

1. If it is hot and humid, then it is raining

$(P \wedge Q) \rightarrow R$

2. If it is humid then it is hot

$Q \rightarrow P$

3. It is humid

Q

Propositional logic: Semantics

- The meaning of a sentence determines by its interpretation
- A sentence is interpreted in terms of **models**, or **possible worlds**.
- These are formal structures that specify a truth value for **each sentence** in a consistent manner.
- m is a model of a sentence α or m satisfies α if α is true in m
- $M(\alpha)$ is the set of all models of α
- A model for KB is a possible world- an assignment of truth values to propositional symbols that makes each sentence in KB is true.
- Possible worlds \sim models
 - Possible worlds: potentially real environments
 - Models: mathematical abstractions that establish the truth or falsity of every sentence
- Example:
 - $x + y = 4$, where $x = \#men$, $y = \#women$
Possible models = all possible assignments of integers to x and y .

Truth tables for connectives

P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$	$P \Leftrightarrow Q$
<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>true</i>
<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>
<i>true</i>	<i>false</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>
<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>

Propositional logic: Semantics

Each model/world specifies true or false for each proposition symbol

E.g. $P_{1,2}$ $P_{2,2}$ $P_{3,1}$
false true false

With these symbols, 8 possible models, can be enumerated automatically.

Rules for evaluating truth with respect to a model m :

$\neg S$ is true iff S is false

$S_1 \wedge S_2$ is true iff S_1 is true and S_2 is true

$S_1 \vee S_2$ is true iff S_1 is true or S_2 is true

$S_1 \Rightarrow S_2$ is true iff S_1 is false or S_2 is true
i.e., is false iff S_1 is true and S_2 is false

$S_1 \Leftrightarrow S_2$ is true iff $S_1 \Rightarrow S_2$ is true and $S_2 \Rightarrow S_1$ is true

Simple recursive process evaluates an arbitrary sentence, e.g.,

$$\neg P_{1,2} \wedge (P_{2,2} \vee P_{3,1}) = \text{true} \wedge (\text{true} \vee \text{false}) = \text{true} \wedge \text{true} = \text{true}$$

Entailment

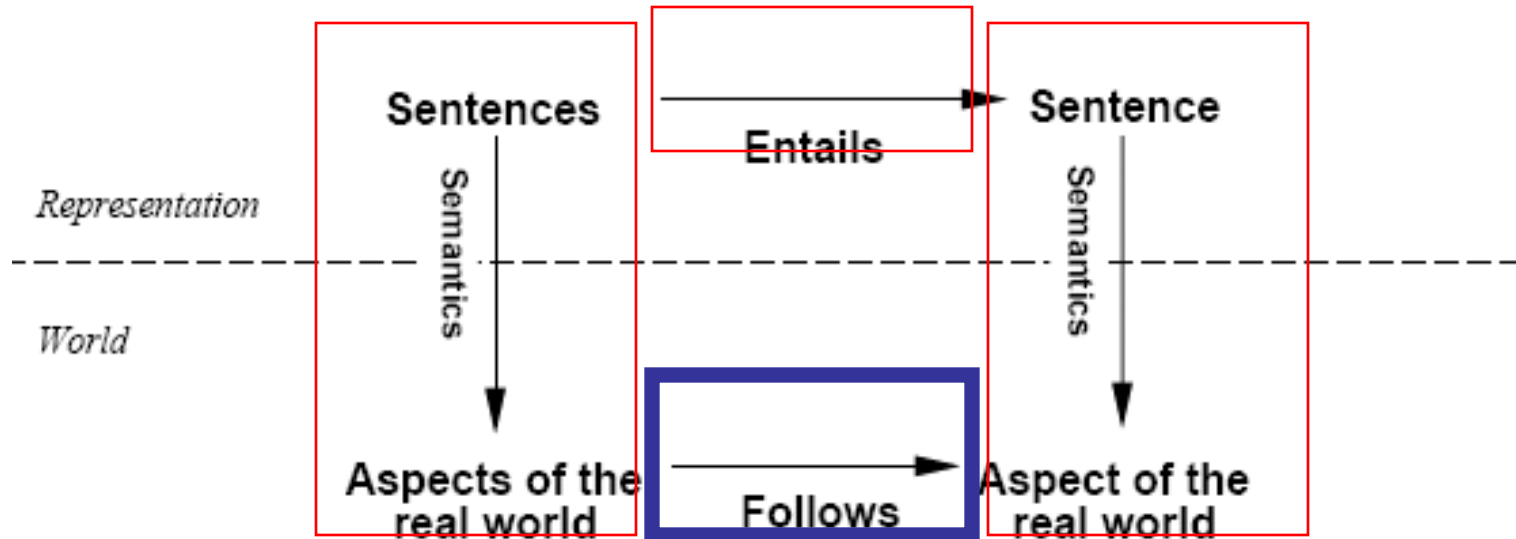
- One sentence follows logically from another

$$\alpha \models \beta$$

α entails sentence β if and only if β is true in all worlds where α is true.

- Entailment is a relationship between sentences that is based on semantics
- Directly related to logical inference

Schematic perspective



If KB is true in the real world, then any sentence α derived from KB by a sound inference procedure is also true in the real world.

Inference Procedures

- *Logical inference creates new sentences that logically follow from a set of sentences (KB)*
- $KB \vdash_i a$ = sentence a can be derived from KB by procedure i
- **Soundness**: i is sound if whenever $KB \vdash_i a$, it is also true that $KB \models a$ (no wrong inferences, but maybe not all inferences)
- **Completeness**: i is complete if whenever $KB \models a$, it is also true that $KB \vdash_i a$ (all inferences can be made, but maybe some wrong extra ones as well)

Inference by enumeration

Model checking approach

- We want to see if a is entailed by KB
- Enumeration of all models
- Check that a is true in every model in which KB is true
- For n symbols, time complexity is $O(2^n)$...
- We need a smarter way to do inference!
- In particular, we are going to infer new logical sentences from the data-base and see if they match a query.

Inference by enumeration Example

KB : $P \wedge Q \rightarrow R$ (R1)
 $Q \rightarrow P$ (R2)

5 models

$R4 : \neg Q \vee R$

Check R4 is true for all the model

P	Q	R	R1	R2	KB	R4
F	F	F	T	T	T ✓	T
F	F	T	T	T	T ✓	T
F	T	F	T	F	NO	
F	T	T	T	F	NO	
T	F	F	T	T	T ✓	T
T	F	T	T	T	T ✓	T
T	T	F	F	T	NO	
T	T	T	T	T	T ✓	T

New KB:

$P \wedge Q \rightarrow R$

$Q \rightarrow P$

$\neg Q \vee R$

Logical equivalence

- To manipulate logical sentences we need some rewrite rules.
- Two sentences are **logically equivalent** iff they are true in same models: $\alpha \equiv \beta$ iff $\alpha \models \beta$ and $\beta \models \alpha$

$$(\alpha \wedge \beta) \equiv (\beta \wedge \alpha) \quad \text{commutativity of } \wedge$$

$$(\alpha \vee \beta) \equiv (\beta \vee \alpha) \quad \text{commutativity of } \vee$$

$$((\alpha \wedge \beta) \wedge \gamma) \equiv (\alpha \wedge (\beta \wedge \gamma)) \quad \text{associativity of } \wedge$$

$$((\alpha \vee \beta) \vee \gamma) \equiv (\alpha \vee (\beta \vee \gamma)) \quad \text{associativity of } \vee$$

$$\neg(\neg\alpha) \equiv \alpha \quad \text{double-negation elimination}$$

$$(\alpha \Rightarrow \beta) \equiv (\neg\beta \Rightarrow \neg\alpha) \quad \text{contraposition}$$

$$(\alpha \Rightarrow \beta) \equiv (\neg\alpha \vee \beta) \quad \text{implication elimination}$$

$$(\alpha \Leftrightarrow \beta) \equiv ((\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)) \quad \text{biconditional elimination}$$

$$\neg(\alpha \wedge \beta) \equiv (\neg\alpha \vee \neg\beta) \quad \text{de Morgan}$$

$$\neg(\alpha \vee \beta) \equiv (\neg\alpha \wedge \neg\beta) \quad \text{de Morgan}$$

$$(\alpha \wedge (\beta \vee \gamma)) \equiv ((\alpha \wedge \beta) \vee (\alpha \wedge \gamma)) \quad \text{distributivity of } \wedge \text{ over } \vee$$

$$(\alpha \vee (\beta \wedge \gamma)) \equiv ((\alpha \vee \beta) \wedge (\alpha \vee \gamma)) \quad \text{distributivity of } \vee \text{ over } \wedge$$

Validity and satisfiability

A sentence is **valid** or **tautology** if it is true in **all** models,

e.g., *True*, $A \vee \neg A$, $A \Rightarrow A$, $(A \wedge (A \Rightarrow B)) \Rightarrow B$

Validity is connected to inference via the **Deduction Theorem**:

$KB \models \alpha$ if and only if $(KB \Rightarrow \alpha)$ is valid

A sentence is **satisfiable** if it is true in **some** model

e.g., $A \vee B$, C

A sentence is **unsatisfiable** if it is false in **all** models

e.g., $A \wedge \neg A$

Satisfiability is connected to inference via the following:


$KB \models \alpha$ if and only if $(KB \wedge \neg \alpha)$ is unsatisfiable

(there is no model for which $KB = \text{true}$ and α is false)

Normal Form

We like to prove: $KB \models \alpha$
equivalent to : $KB \wedge \neg \alpha$ unsatisfiable

We first rewrite $KB \wedge \neg \alpha$ into conjunctive normal form (CNF).



- Any KB can be converted into CNF

Example: Conversion to CNF

$$B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$$

1. Eliminate \Leftrightarrow , replacing $\alpha \Leftrightarrow \beta$ with $(\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)$.
 $(B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})) \wedge ((P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1})$
2. Eliminate \Rightarrow , replacing $\alpha \Rightarrow \beta$ with $\neg \alpha \vee \beta$.
 $(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg(P_{1,2} \vee P_{2,1}) \vee B_{1,1})$
3. Move \neg inwards using de Morgan's rules and double-negation: $\neg(\alpha \vee \beta) = \neg \alpha \wedge \neg \beta$
 $(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge ((\neg P_{1,2} \wedge \neg P_{2,1}) \vee B_{1,1})$
4. Apply distributive law (\wedge over \vee) and flatten:
 $(\neg B_{1,1} \vee P_{1,2} \vee P_{2,1}) \wedge (\neg P_{1,2} \vee B_{1,1}) \wedge (\neg P_{2,1} \vee B_{1,1})$
A conjunction of three clauses

Sound rules of inference

- Here are some examples of sound rules of inference
- Each can be shown to be sound using a truth table

<u>RULE</u>	<u>PREMISE</u>	<u>CONCLUSION</u>
Modus Ponens	$A, A \rightarrow B$	B
And Introduction	A, B	$A \wedge B$
And Elimination	$A \wedge B$	A
Double Negation	$\neg\neg A$	A
Unit Resolution	$A \vee B, \neg B$	A
Resolution	$A \vee B, \neg B \vee C$	$A \vee C$

Resolution

- **Resolution:** inference rule for CNF: **sound and complete!**

$$(A \vee B \vee C)$$

$$(\neg A)$$

“If A or B or C is true, but not A, then B or C must be true.”

$$\therefore (B \vee C)$$

$$(A \vee B \vee C)$$

$$(\neg A \vee D \vee E)$$

“If A is false then B or C must be true, or if A is true then D or E must be true, hence since A is either true or false, B or C or D or E must be true.”

$$\therefore (B \vee C \vee D \vee E)$$

$$(A \vee B)$$

$$(\neg A \vee B)$$

Simplification

$$\therefore (B \vee B) \equiv B$$

Resolution

- **Resolution** is a valid inference rule producing a new clause implied by two clauses containing *complementary literals*
 - A literal is an atomic symbol or its negation, i.e., P , $\sim P$

Resolution

- A KB is actually a set of sentences all of which are true, i.e., a conjunction of sentences.
- To use resolution, put KB into *conjunctive normal form* (CNF), where each sentence written as a disjunction of (one or more) literals

Example

- KB: $[P \rightarrow Q, Q \rightarrow R \wedge S]$
- KB in CNF: $[\sim P \vee Q, \sim Q \vee R, \sim Q \vee S]$
- Resolve KB(1) and KB(2) producing: $\sim P \vee R$ (i.e., $P \rightarrow R$)
- Resolve KB(1) and KB(3) producing: $\sim P \vee S$ (i.e., $P \rightarrow S$)
- New KB: $[\sim P \vee Q, \sim Q \vee \sim R \vee \sim S, \sim P \vee R, \sim P \vee S]$

Resolution Algorithm

- The resolution algorithm tries to prove: $KB \models \alpha$ equivalent to $KB \wedge \neg\alpha$ unsatisfiable
- $KB \wedge \neg\alpha$ is converted to CNF
- The resolution rule is applied to the resulting clauses.
- Each pair that contains complementary literals is resolved to produce the new clause, which is added if it is not already present.
- Continue until one of two things happens:
 1. There is no new clause that can be added, in which case KB does not entail α (no contradiction; there is a model that satisfies the sentence $KB \wedge \neg\alpha$ (non-trivial))
 2. Two clauses resolve to yield the empty clause, \square , in which case KB entails α . The empty clause represents a contradiction is to observe that it arises only from resolving two complementary unit clauses

Resolution example

- $KB = (B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge \neg B_{1,1}$
- $\alpha = \neg P_{1,2}$

$KB \wedge \neg \alpha$

