



Chapter 7: Entity-Relationship Model

Edited by Radhika Sukapuram

Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for 7.1 conditions on re-use

©Silberschatz, Korth and Sudarshan



Covered so far

- Assumed a database schema
- Designed queries, constraints and updates
- Understood security schemes

But how to arrive at a database schema ?

For complex real world applications, one cannot directly arrive at a database schema



Design Phases

- Understand the data needs of the prospective database users by talking to users and domain experts
 - ▶ Output: textual descriptions of user requirements
- Choose a data model
 - Apply the concepts of the chosen data model, translate requirements into a conceptual schema of the database.
 - Entity-Relationship model is typically used
 - What are the entities ? What are there relationships between them ?
 - Output: Entity-Relationship diagram and functional requirements
 - Review:
 - ▶ to find out if requirements are met
 - ▶ No conflicts
 - ▶ No redundant features



Design Phases cont.

- Functional requirements describe
 - The kinds of operations – transactions – that will be performed on the data
- Implementation
 - Logical design
 - ▶ map E-R diagram to a relational model
 - ▶ Business decision – What attributes should we record in the database?
 - ▶ Computer Science decision – What relation schemas should we have and how should the attributes be distributed among the various relation schemas?
 - Physical design – file organisation, index structures
- Changes after implementation
 - Changes to logical schema are hard – affect application code
 - Do logical design with care!



Design Approaches

- Entity Relationship Model
 - Models an enterprise as a collection of *entities* and *relationships*
 - ▶ Entity: a “thing” or “object” in the enterprise that is distinguishable from other objects
 - Described by a set of *attributes*
 - ▶ Relationship: an association among several entities
 - Represented diagrammatically by an *entity-relationship diagram*:
- Normalization Theory (later)
 - Formalize what designs are bad, and test for them



Outline of the ER Model



ER model -- Database Modeling

- ❑ The ER data model - to facilitate database design
 - ❑ by allowing specification of an **enterprise schema**
 - ❑ The enterprise schema represents the overall logical structure of a database.
 - ❑ Useful in mapping the meanings and interactions of real-world enterprises onto a conceptual schema.
- ❑ The ER data model employs three basic concepts:
 - ❑ entity sets,
 - ❑ relationship sets,
 - ❑ attributes.
- ❑ ER diagram is a graphical representation



Entity Sets

- An **entity** is an object that exists and is distinguishable from other objects.
 - Example: a specific person, company, event, an instructor
- An **entity set** is a set of entities of the same type that share the same properties.
 - Example: set of all persons, companies, instructors
- An entity is represented by a set of attributes
 - descriptive properties possessed by all members of an entity set.
 - Example:
 $instructor = (ID, name, street, city, salary)$
 $course = (course_id, title, credits)$
- A minimal subset of the attributes form a **primary key** of the entity set; i.e., uniquely identifying each member of the set.



Entity Sets -- *instructor* and *student*

instructor_ID instructor_name

| | |
|-------|------------|
| 76766 | Crick |
| 45565 | Katz |
| 10101 | Srinivasan |
| 98345 | Kim |
| 76543 | Singh |
| 22222 | Einstein |

instructor

student-ID student_name

| | |
|-------|---------|
| 98988 | Tanaka |
| 12345 | Shankar |
| 00128 | Zhang |
| 76543 | Brown |
| 76653 | Aoi |
| 23121 | Chavez |
| 44553 | Peltier |

student



Relationship Sets

- A **relationship** is an association among several entities

Example:

| | | |
|-----------------------|------------------|---------------------------|
| 44553 (Peltier) | <u>advisor</u> | 22222 (<u>Einstein</u>) |
| <i>student</i> entity | relationship set | <i>instructor</i> entity |

- A **relationship set** is a mathematical relation among $n \geq 2$ entities, each taken from entity sets

$$\{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$

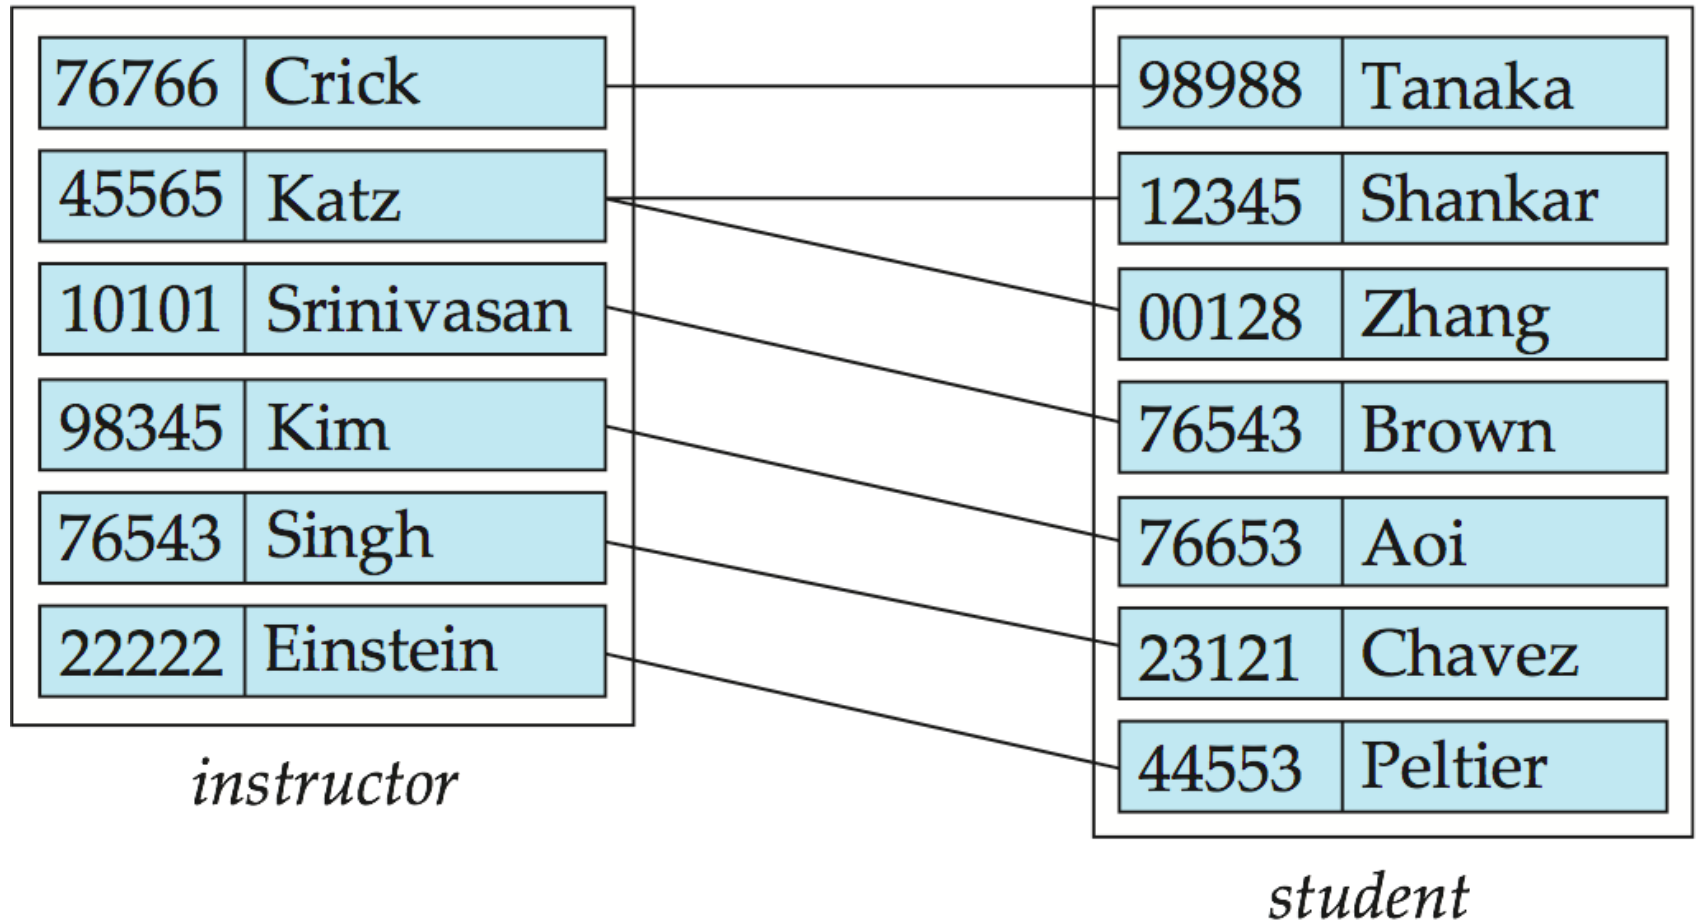
where (e_1, e_2, \dots, e_n) is a relationship

- Example:

$$(44553, 22222) \in \text{advisor}$$



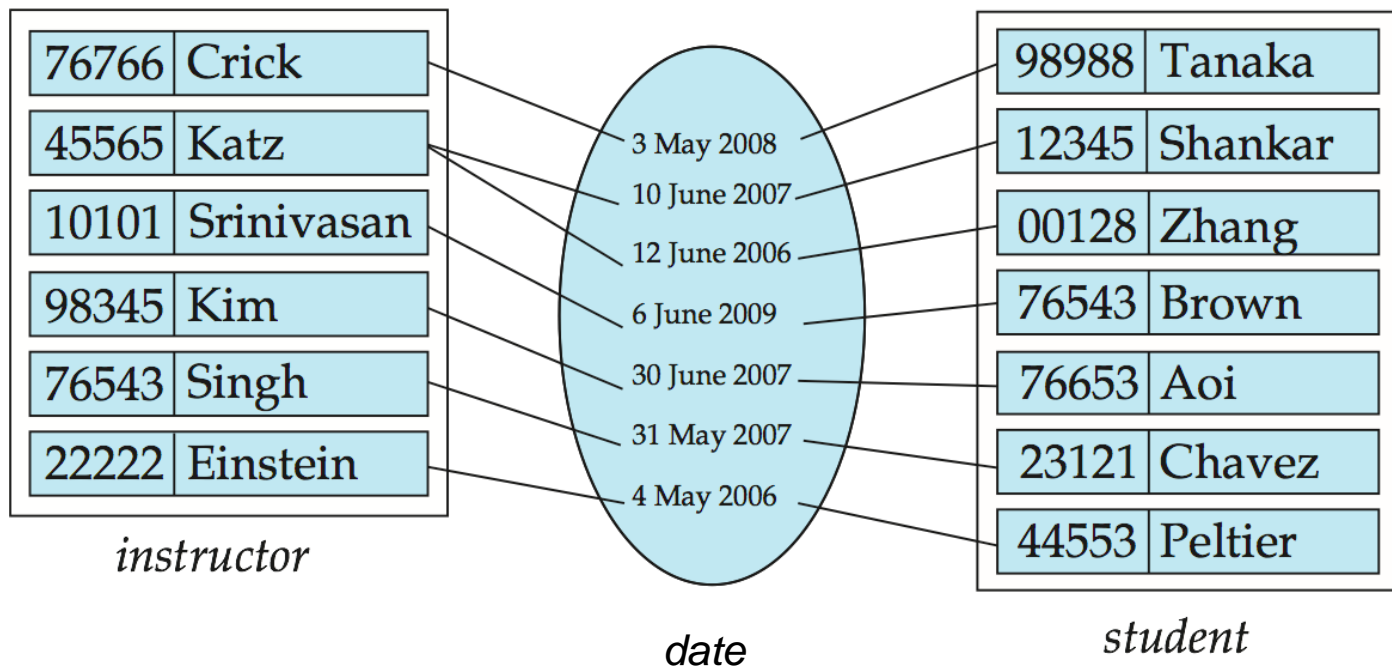
Relationship Set *advisor*





Relationship Sets (Cont.)

- An attribute can also be associated with a relationship set.
- For instance, the *advisor* relationship set between entity sets *instructor* and *student* may have the attribute *date*
- *date* tracks when the student started being associated with the advisor





Degree of a Relationship Set

- binary relationship
 - involves two entity sets (or degree two).
- Relationships between more than two entity sets are rare. Most relationships are binary. (More on this later.)
 - ▶ Example: *students* work on research *projects* under the guidance of an *instructor*.
 - ▶ relationship *proj_guide* is a ternary relationship between *instructor*, *student*, and *project*



Roles

- Entity sets of a relationship need not be distinct
 - If distinct, the function an entity set plays is implicit
 - Otherwise each occurrence of an entity set plays a “role” in the relationship
 - The same entity set participates in a relationship set more than once, in **different roles**
 - Called recursive relationship set
- One *course* is a pre-requisite for another *course*
 - Use a relationship set called *prereq*
 - ▶ Modelled by ordered pairs of *course* entities
 - ▶ (*course_role*, *prereq_role*) – *prereq_role* is the pre-requisite
 - ▶ (*prereq_role*, *course_role*) pairs are excluded
 - ▶ *prereq* is a relationship between *course_role* and *prereq_role*





Mapping Cardinality Constraints

- Mapping cardinalities or cardinality ratios:
 - Express the number of entities to which another entity can be associated via a relationship set.
- Most useful in describing binary relationship sets.
- For a binary relationship set the mapping cardinality must be one of the following types:
 - One to one
 - One to many
 - Many to one
 - Many to many




Mapping cardinalities


- One-to-one
 - An entity in A is associated with at most one entity in B and an entity in B is associated with at most one entity in A
- One-to-many
 - An entity in A is associated with zero or more entities in B. An entity in B is associated with at most one entity in A



Mapping Cardinalities



One to one




One to many

Note: Some entities in A and B may not be mapped to any entities in the other set




Mapping Cardinalities



(a)

Many to
one



(b)

Many to many

Note: Some elements in A and B may not be mapped to any elements in the other set



Participation Constraints

- Participation of an entity set E in a relationship R is total if every entity in E participates in at least one relationship in R
 - Every *student* entity must be related to at least one *instructor* through the *advisor* relationship
- If only some entities of an entity set E participate in a relationship R , the participation is partial
 - Only some of the *instructor* entities are related to *student* entities through the relationship *advisor*



Keys for relationship sets

- Superkeys, candidate keys and primary key apply to entity sets
- R is a relationship set involving entities E_1, E_2, \dots, E_n
- $primarykey(E_1)$ denotes the set of attributes that form the primary key of E_1
- Assume that attribute names of all primary keys are unique
 - Suppose R has no attributes associated with it
 - ▶ The set of attributes
 $primarykey(E_1) \cup primarykey(E_2) \cup \dots \cup primarykey(E_n)$ describes a relationship in R
 - Suppose R has attributes a_1, a_2, \dots, a_m associated with it
 - ▶ The set of attributes
 $primarykey(E_1) \cup primarykey(E_2) \cup \dots \cup primarykey(E_n) \cup \{a_1, a_2, \dots, a_m\}$ describes a relationship in R



Keys for relationship sets cont.

- Whether R has attributes or not,
 - $primarykey(E_1) \cup primarykey(E_2) \cup \dots \cup primarykey(E_n)$ form a superkey for R
- Attribute names of primary keys across entities are not unique
 - Attributes are renamed : <name of entity set> . <name of attribute>
- An entity set participates more than once in a relationship
 - Attributes are renamed : <role name> . <name of attribute>
 - Example: courses and pre-requisites





Primary keys for a relationship set

- Structure of primary key of a relationship set depends on mapping cardinality
- Suppose *advisor* is a relationship set between *instructor* and *student*
 - many-to-many : union of primary keys of *instructor* and *student*
 - many-to-one from *student* to *instructor* : each *student* can be advised by at most one *instructor* : primary key of *student*
 - many-to-one from *instructor* to *student* : each *instructor* can advise at most one *student* : primary key of *instructor*
 - one-to-one: either candidate key can be used as the primary key



Primary keys: non-binary relationship sets

- No cardinality constraints:

$primarykey(E_1) \cup primarykey(E_2) \cup \dots \cup primarykey(E_n)$ form a superkey for R and is the only candidate key

- With cardinality constraints: specified later