



Chapter 7: Entity-Relationship Model

Edited by Radhika Sukapuram

Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for 7.1 conditions on re-use

©Silberschatz, Korth and Sudarshan



Reduction to Relation Schemas



Reduction to Relation Schemas

- Entity sets and relationship sets can be expressed uniformly as *relation schemas*.
- A database which conforms to an E-R diagram can be represented by a collection of schemas.
- For each entity set and relationship set there is
 - a unique schema that is assigned the name of the corresponding entity set or relationship set.
- Each schema has a number of columns (generally corresponding to attributes), which have unique names.



Representing Entity Sets


- A strong entity set reduces to a schema with the same attributes

student(ID, name, tot_cred)

and the same primary keys

- A weak entity set becomes a table that includes a column for the primary key of the identifying strong entity set


section (course id, sec id, sem, year)





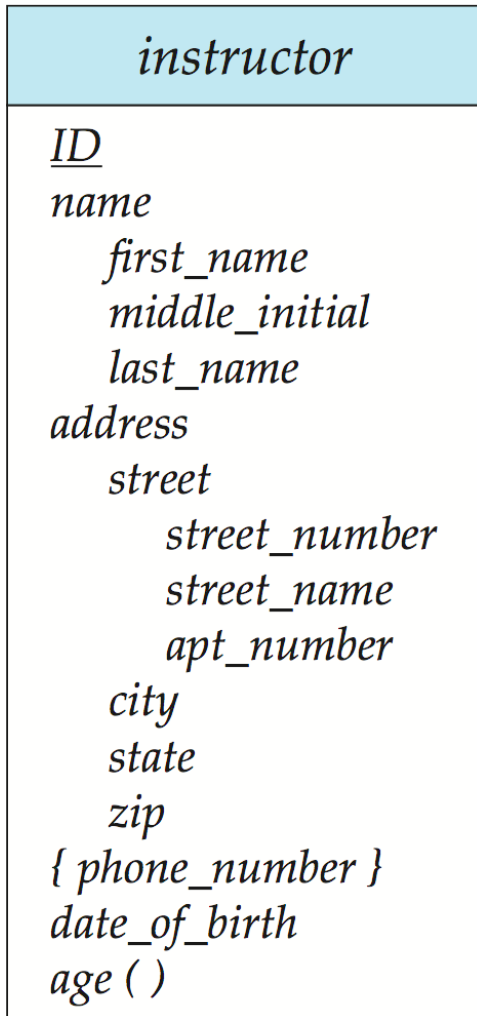
Representing entity sets cont.

- A foreign key constraint on *section* is created
 - *course_id* of *section* references the primary key *course_id* of *course*
 - with on delete cascade





Representation of Entity Sets with Composite Attributes



- Composite attributes are flattened out by creating a separate attribute for each component attribute
 - the schema corresponding to the entity set *instructor* has three attributes *name_first_name* , *name_middle_initial* and *name_last_name*
 - ▶ Prefix omitted if there is no ambiguity (*name_first_name* could be *first_name*)
- Ignoring multivalued attributes, extended *instructor* schema is
 - *instructor*(*ID*,
 first_name, *middle_initial*, *last_name*,
 street_number, *street_name*,
 apt_number, *city*, *state*, *zip_code*,
 date_of_birth)



Representation of Entity Sets with Multivalued Attributes

- Example: Multivalued attribute *phone_number* of *instructor* is represented by a schema:
$$inst_phone = (\underline{ID}, \underline{phone_number})$$
 - Why is *phone_number* a part of the primary key ?
- A multivalued attribute *M* of an entity *E* is represented by a separate schema *EM*
- Schema *EM* has attributes corresponding to the primary key of *E* and an attribute corresponding to multivalued attribute *M*



Representation of Entity Sets with Multivalued Attributes cont.

- Each value of the multivalued attribute maps to a separate tuple of the relation on schema *EM*
 - For example, an *instructor* entity with primary key 22222 and phone numbers 456-7890 and 123-4567 maps to two tuples:
(22222, 456-7890) and (22222, 123-4567)
- Create a foreign key constraint from EM to E
 - primary key derived from E of EM referencing E
 - *id* of *inst_phone* references *instructor*
- Exercise: What if E has only two attributes : one simple attribute which is a primary key and a multivalued attribute ?




Representing Relationship Sets

- A many-to-many relationship set is represented as
 - a schema with attributes for the primary keys of the two participating entity sets
 - and any descriptive attributes of the relationship set.
- Example: schema for relationship set *advisor*

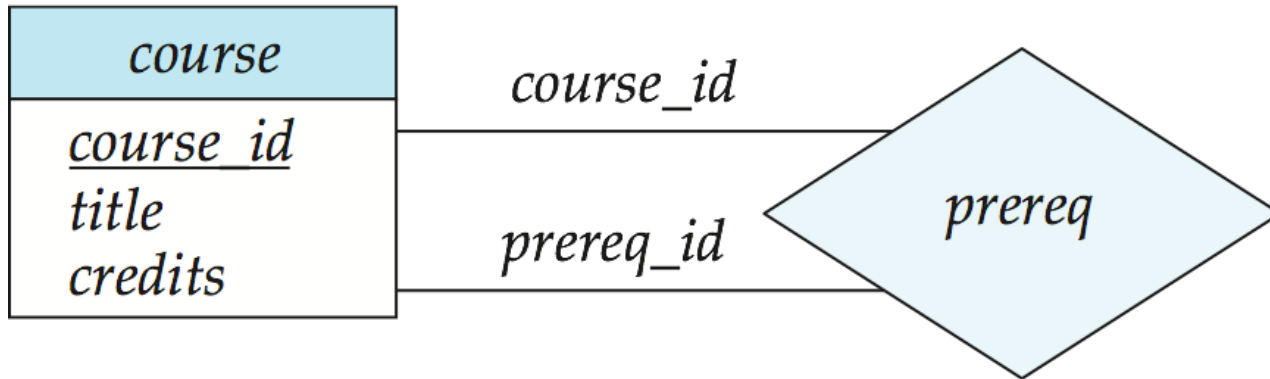
$advisor = (\underline{s_id}, \underline{i_id})$

- Two foreign keys, with *s_id* referencing student and *i_id* referencing *instructor*





Representing Relationship Sets




prereq (*course id*, *prereq id*)



Non-binary Relationship Sets


- ❑ Most relationship sets are binary
- ❑ There are occasions when it is more convenient to represent relationships as non-binary.
- ❑ E-R Diagram with a Ternary Relationship





Cardinality Constraints on Ternary Relationships


- We allow at most one arrow out of a ternary (or greater degree) relationship to indicate a cardinality constraint
- For example, an arrow from *proj_guide* to *instructor* indicates each student has at most one guide for a project





Cardinality Constraints on Ternary Relationships cont.

- If there is more than one arrow, there are two ways of defining the meaning.
 - For example, a ternary relationship R between A , B and C with arrows to B and C could mean
 1. Each A entity is associated with a unique entity from B and C or
 2. Each pair of entities from (A, B) is associated with a unique C entity, and each pair (A, C) is associated with a unique B
 - Each alternative has been used in different formalisms
 - To avoid confusion we outlaw more than one arrow





Representing Relationship Sets cont.

- Exercise: What are the primary keys for schemes translated from binary relationships with various cardinality mappings ?
- n-ary relationship without any arrows on the edges
 - Union of the primary key attributes of the participating relationships
- n-ary relationship with an arrow on one of its edges (only 1 arrow is allowed)
 - Union of primary key of the entity sets not on the arrow side of the relationship



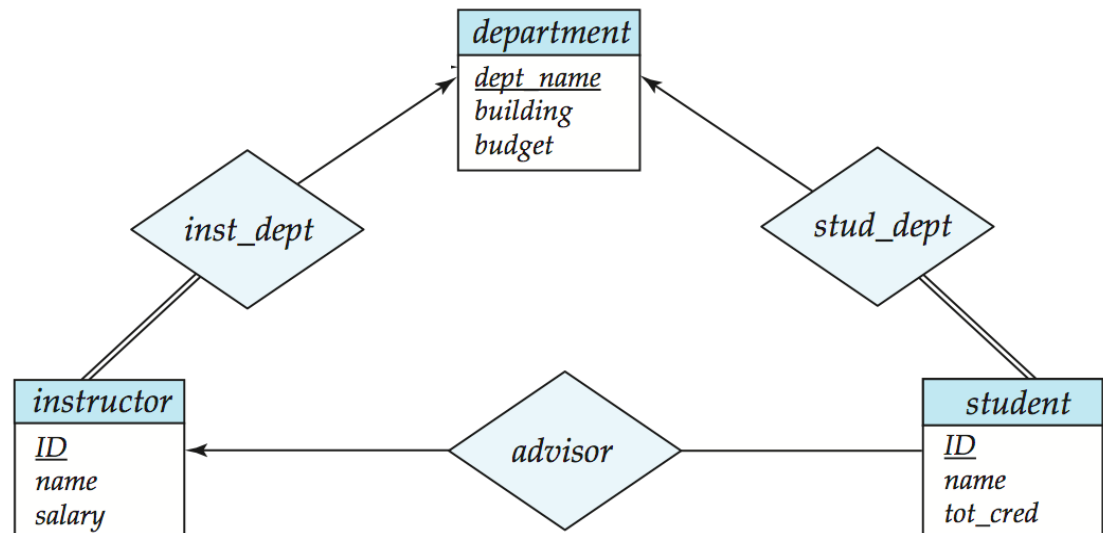
Representing Relationship Sets cont.

- Foreign keys must also be created from relation R
 - For each entity set E_i related to R
 - Create a foreign key from R
 - With the attributes of R that were derived from the primary key attributes of E_i
 - Referencing the primary key of the relation representing E_i
- Exercise : Are foreign keys always created from relations created from relationship sets to the relations created from the entity sets they were derived from ?



Combination of Schemas

- Instead of creating a schema for relationship set *inst_dept*
 - add an attribute *dept_name* to the schema arising from entity set *instructor*
- Many-to-one and one-to-many relationship sets that are total on the many-side
 - can be represented by adding extra attributes to the “many” side
 - containing the primary key of the “one” side
 - Add a foreign key to the “many” side referencing the “one” side
 - Primary key: the primary key of entity set into whose schema the merge took place






Combination of Schemas (Cont.)

- For one-to-one relationship sets, either side can be chosen to act as the “many” side
 - That is, an extra attribute can be added to either of the tables corresponding to the two entity sets
- If participation is *partial* on the “many” side (*instructor*),
 - replacing a schema by an extra attribute in the schema corresponding to the “many” side , containing the primary key of the “one” side (*dept_name*) could result in null values (for *dept_name*)



Redundancy of Schemas

- The schema corresponding to a relationship set linking a weak entity set to its identifying strong entity set is redundant.
- Example: The *section* schema already contains the attributes that would appear in the *sec_course* schema



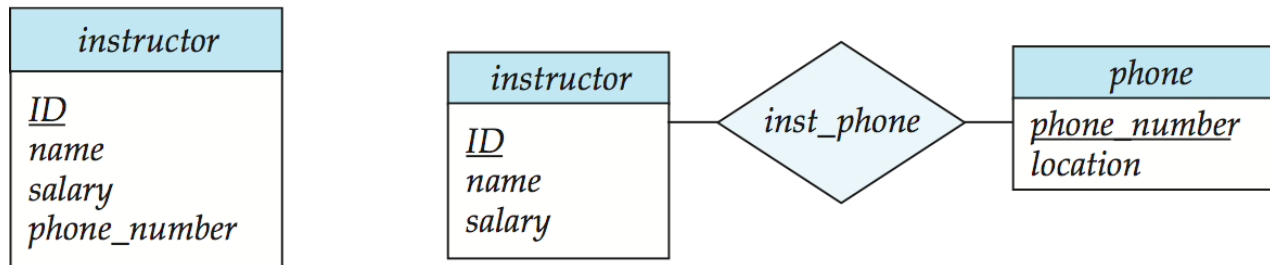


Design Issues



Entities vs. Attributes

- Use of entity sets vs. attributes



- Use of phone as an entity allows extra information about phone numbers (plus multiple phone numbers)
- An entity is more general than an attribute
- Use an entity when generality is more useful



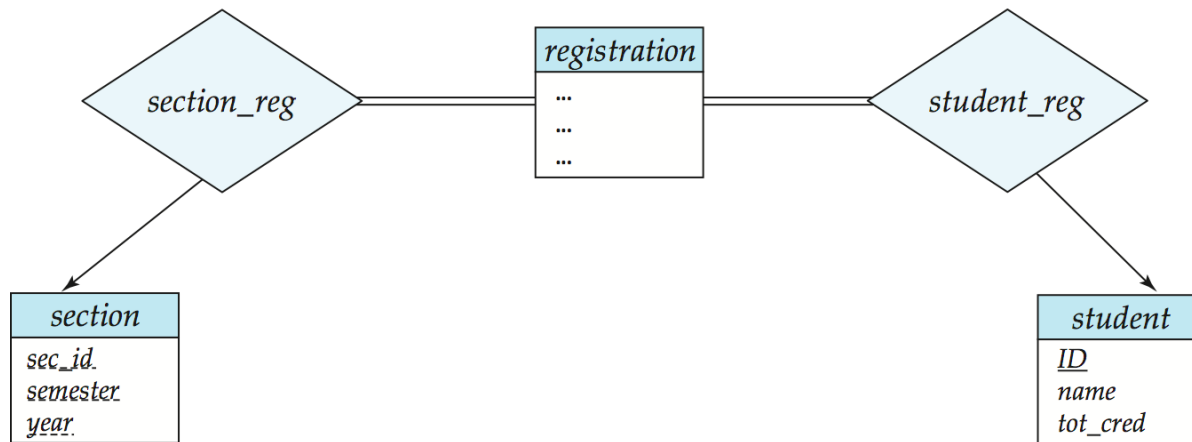
Entities vs attributes

- Common mistakes
 - Instead of using a relationship (*advisor*)
 - ▶ using the primary key of an entity set (*id* of *student*) as an attribute of another entity set (*instructor*)
 - ▶ *id* of a *student* must not be an attribute of *instructor* even if the instructor advises only one student
 - ▶ Instead, use an *advisor* relationship
 - Makes the connection explicit
 - A better representation of reality
 - Primary key attributes of entity sets are made attributes of the relationship set
 - ▶ *id* of a *student* and *id* of *instructor* must not be attributes of *advisor*
 - Because the primary key attributes are implicit in the relationship!



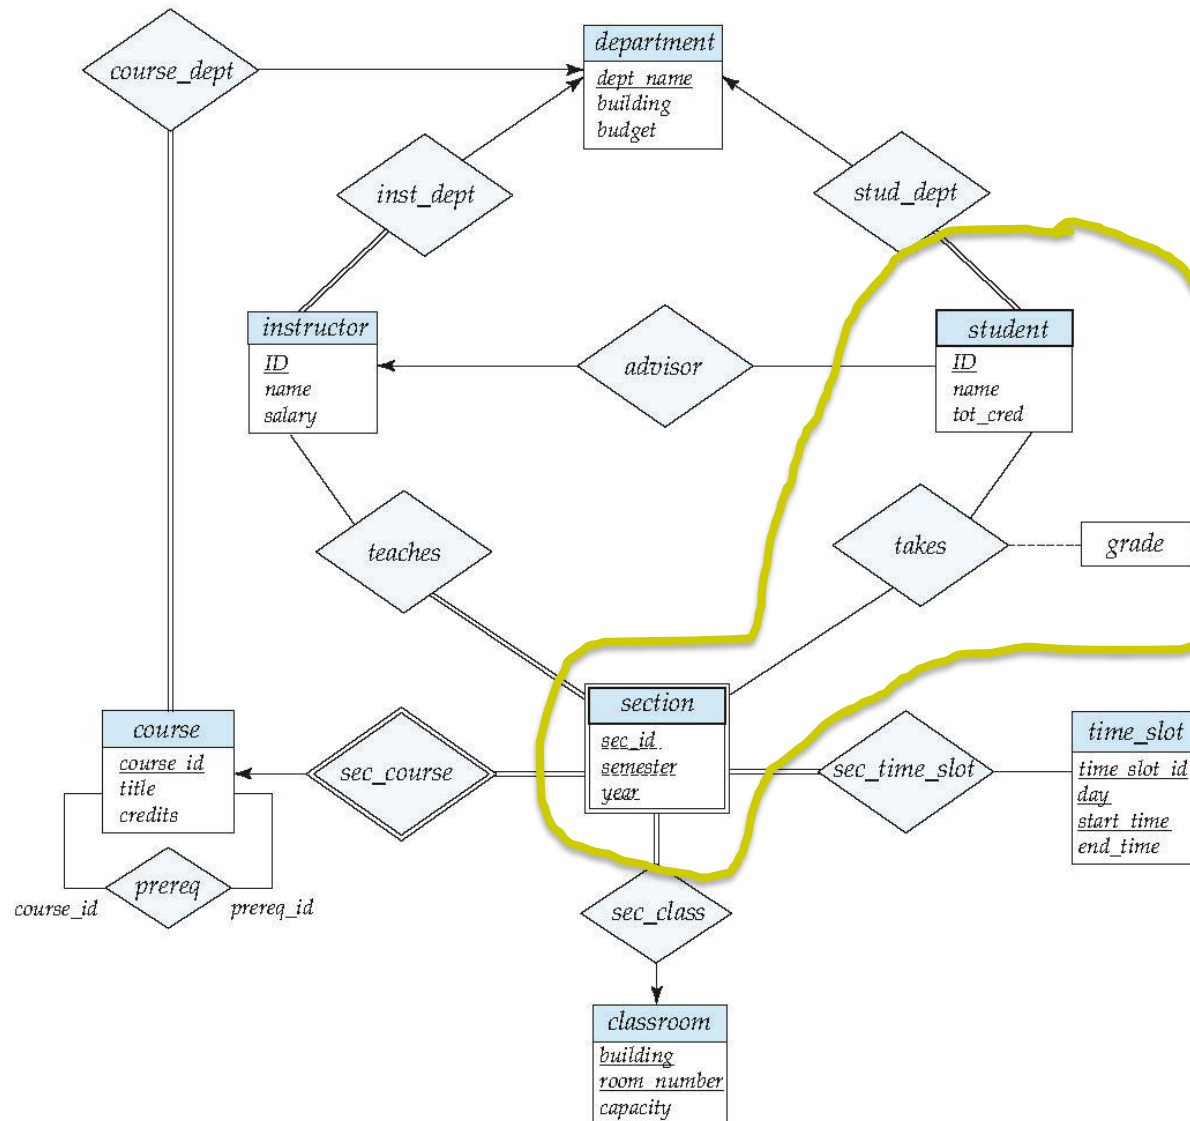
Entities vs. Relationship sets

- **Use of entity sets vs. relationship sets**
 - A student takes a section of a course
 - Instead, create a course registration record for each course that the student takes
- **Possible guidelines**
 - If more attributes are associated with registration, it is an entity in its own right
 - designate a relationship set to describe an action that occurs between entities





E-R Diagram for a University Enterprise





Entities vs. Relationship sets

□ Placement of relationship attributes

- For example, attribute *date* (date on which an instructor starts advising a student) as attribute of *advisor* or as attribute of *student*
- One-to-many from *instructor* - *student*: *date* may be associated with *student* (the entity on the many side)
- One-to-one: either of the participating entities
- Also depends on the characteristics of the enterprise being modelled
 - ▶ *date* may be associated with *advisor* to indicate that it is the date on which the advisor was associated with the student and not the student's date of entry into the university
- Many-to-many: associated with the relationship set




Binary Vs. Non-Binary Relationships

- It is possible to replace any non-binary (n -ary, for $n > 2$) relationship set by a number of distinct binary relationship sets
 - Why ? some data models support only binary relationship sets
- But an n -ary relationship set shows more clearly that several entities participate in a single relationship.
- Some relationships that appear to be non-binary may be better represented using binary relationships
 - For example, a ternary relationship *parents*, relating a child to his/her father and mother, is best replaced by two binary relationships, *father* and *mother*
 - ▶ Using two binary relationships allows partial information (e.g., only one parent being known)
 - ▶ Otherwise null values will need to be used
 - But there are some relationships that are naturally non-binary
 - ▶ Example: *proj_guide*




Converting Non-Binary Relationships to Binary Form

- In general, any non-binary relationship can be represented using binary relationships by creating an artificial entity set.
 - Replace R between entity sets A , B and C by an entity set E , and three relationship sets:
 1. R_A , relating E and A
 2. R_B , relating E and B
 3. R_C , relating E and C
 - Create an identifying attribute for E and add any attributes of R to E
 - For each relationship (a_i, b_i, c_i) in R , create
 1. a new entity e_i in the entity set E
 2. add (e_i, a_i) to R_A
 3. add (e_i, b_i) to R_B
 4. add (e_i, c_i) to R_C



(a)



(b)



Converting Non-Binary Relationships (Cont.)

- Also need to translate constraints
 - Translating all constraints may not be possible
 - ▶ Each pair of entities from A and B may be associated with at most one entity C – how to translate this ?



Restricting to binary relationship sets

- Restricting to binary relationship sets is not always desirable:
 - Increases space and complexity
 - Does not reflect reality intuitively
 - Translating all constraints may not be possible