



Chapter 8: Relational Database Design

Database System Concepts, 6th Ed.

Edited by Radhika Sukapuram

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Testing for Dependency Preservation

- To check if a dependency $\alpha \rightarrow \beta$ is preserved in a decomposition of R into R_1, R_2, \dots, R_n we apply the following test (with attribute closure done with respect to F)
 - $result = \alpha$
repeat
 - for each** R_i in the decomposition
 - $t = (result \cap R_i)^+ \cap R_i$
 - $result = result \cup t$**until** ($result$ does not change)
 - If $result$ contains all attributes in β , then the functional dependency $\alpha \rightarrow \beta$ is preserved.
- We apply the test on all dependencies in F to check if a decomposition is dependency preserving
- This procedure takes polynomial time, instead of the exponential time required to compute F^+ and $(F_1 \cup F_2 \cup \dots \cup F_n)^+$



Example

- $R = (A, B, C, D, G)$
- $F = \{A \rightarrow B, B \rightarrow C, A \rightarrow D, D \rightarrow G\}$
- Case: $R_1(A, B)$, $R_2(B, C, D)$, $R_3(D, G)$, Whether $A \rightarrow G$?





Testing for BCNF

- To check if a non-trivial dependency $\alpha \rightarrow \beta$ causes a violation of BCNF
 1. compute α^+ (the attribute closure of α), and
 2. verify that it includes all attributes of R , that is, it is a superkey of R .

- **Simplified test:** To check if a relation schema R is in BCNF,
 - it suffices to check only the dependencies in the given set F for violation of BCNF, rather than checking all dependencies in F^+ .
 - If none of the dependencies in F causes a violation of BCNF, then none of the dependencies in F^+ will cause a violation of BCNF either.



Testing for BCNF

- However, **simplified test using only F is incorrect when testing a relation in a decomposition of R**
 - Consider $R = (A, B, C, D, E)$, with $F = \{ A \rightarrow B, BC \rightarrow D \}$
 - ▶ Decompose R into $R_1 = (A, B)$ and $R_2 = (A, C, D, E)$
 - ▶ Neither of the dependencies in F contain only attributes from (A, C, D, E) so we might be misled into thinking R_2 satisfies BCNF.
 - ▶ In fact, dependency $AC \rightarrow D$ in F^+ shows R_2 is not in BCNF.
 - ▶ So we need to check with all dependencies in F^+ to check if a decomposition is in BCNF.



Easier :Testing Decomposition for BCNF

- To check if a relation R_i in a decomposition of R is in BCNF,
 - Either
 - ▶ Find F^+
 - ▶ Find **restriction** of F to R_i (that is, all FDs in F^+ that contain only attributes from R_i)
 - ▶ test R_i for BCNF with respect to the restriction of F to R_i
 - or **use the original set of dependencies F that holds on R** , but with the following test:
 - ▶ for every set of attributes $\alpha \subseteq R_i$, check that α^+ (the attribute closure of α under F)
 - » either includes no attribute of $R_i - \alpha$,
 - » or includes all attributes of R_i .
 - ▶ If the condition is violated by some α in R_i , the dependency $\alpha \rightarrow (\alpha^+ - \alpha) \cap R_i$ can be shown to hold on R_i , and R_i violates BCNF.



BCNF Decomposition Algorithm

```
result := {R};
done := false;
compute  $F^+$ ;
while (not done) do
    if (there is a schema  $R_i$  in result that is not in BCNF)
        then begin
            let  $\alpha \rightarrow \beta$  be a nontrivial functional dependency that
                holds on  $R_i$  such that  $\alpha \rightarrow R_i$  is not in  $F^+$ ,
                and  $\alpha \cap \beta = \emptyset$ ;
            result := (result -  $R_i$ )  $\cup$  ( $R_i - \beta$ )  $\cup$  ( $\alpha, \beta$ );
        end
    else done := true;
```

Note: each R_i is in BCNF, and decomposition is lossless-join.



Example of BCNF Decomposition

- $R = (A, B, C)$
 $F = \{B \rightarrow C, A \rightarrow B\}$
Key = $\{A\}$
- R is not in BCNF ($B \rightarrow C$ but B is not a superkey)
- $F^+ = \{B \rightarrow C, A \rightarrow B, A \rightarrow C\}$
- $B \rightarrow C$ is an FD such that $B \rightarrow ABC$ is not in F^+ and $B \cap C = \emptyset$
- So remove ABC, add AC and BC
- Decomposition
 - $R_1 = (B, C)$
 - $R_2 = (A, B)$



Example of BCNF Decomposition

- *class* (*course_id*, *title*, *dept_name*, *credits*, *sec_id*, *semester*, *year*, *building*, *room_number*, *capacity*, *time_slot_id*)
- Functional dependencies:
 - *course_id* → *title*, *dept_name*, *credits*
 - *building*, *room_number* → *capacity*
 - *course_id*, *sec_id*, *semester*, *year* → *building*, *room_number*, *time_slot_id*
- A candidate key {*course_id*, *sec_id*, *semester*, *year*}.
- BCNF Decomposition:
 - *course_id* → *title*, *dept_name*, *credits* holds
 - ▶ but *course_id* is not a superkey. *class* is not in BCNF
 - We replace *class* by:
 - ▶ *course*(*course_id*, *title*, *dept_name*, *credits*)
 - ▶ *class-1* (*course_id*, *sec_id*, *semester*, *year*, *building*, *room_number*, *capacity*, *time_slot_id*)



BCNF Decomposition (Cont.)

- *course* is in BCNF
 - How do we know this ?
- *building, room_number* → *capacity* holds on *class-1*
 - but {*building, room_number*} is not a superkey for *class-1*.
 - We replace *class-1* by:
 - ▶ *classroom* (*building, room_number, capacity*)
 - ▶ *section* (*course_id, sec_id, semester, year, building, room_number, time_slot_id*)
- *classroom* and *section* are in BCNF.



BCNF and Dependency Preservation

It is not always possible to get a BCNF decomposition that is dependency preserving

□ $R = (J, K, L)$

$F = \{JK \rightarrow L$
 $L \rightarrow K\}$

Two candidate keys = JK and JL

□ R is not in BCNF

□ Any decomposition of R will fail to preserve

$JK \rightarrow L$

This implies that testing for $JK \rightarrow L$ requires a join



Third Normal Form: Motivation

- There are some situations where
 - BCNF is not dependency preserving, and
 - efficient checking for FD violation on updates is important
- Solution: define a weaker normal form, called Third Normal Form (3NF)
 - Allows some redundancy (with resultant problems; we will see examples later)
 - But functional dependencies can be checked on individual relations without computing a join.
 - There is always a lossless-join, dependency-preserving decomposition into 3NF.



Third Normal Form

- A relation schema R is in **third normal form (3NF)** if for all:

$$\alpha \rightarrow \beta \text{ in } F^+$$

at least one of the following holds:

- $\alpha \rightarrow \beta$ is trivial (i.e., $\beta \subseteq \alpha$)
- α is a superkey for R
- Each attribute A in $\beta - \alpha$ is contained in a candidate key for R .

(**NOTE:** each attribute may be in a different candidate key)

- If a relation is in BCNF it is in 3NF
 - since in BCNF one of the first two conditions above must hold.
- The third condition is a minimal relaxation of BCNF to ensure dependency preservation (will see why later).
- A relation in 3NF may have redundancies due to FDs



3NF Example

- Relation *dept_advisor*.
 - *dept_advisor* (*s_ID*, *i_ID*, *dept_name*)
 $F = \{s_ID, dept_name \rightarrow i_ID, i_ID \rightarrow dept_name\}$
 - Two candidate keys: *s_ID*, *dept_name*, and *i_ID*, *s_ID*
 - *R* is in 3NF
 - ▶ $s_ID, dept_name \rightarrow i_ID$
 - *s_ID*, *dept_name* is a superkey
 - ▶ $i_ID \rightarrow dept_name$
 - *dept_name* is contained in a candidate key



Redundancy in 3NF

- There is some redundancy in this schema
- Example of problems due to redundancy in 3NF

- $R = (J, K, L)$
 $F = \{JK \rightarrow L, L \rightarrow K\}$

J	L	K
j_1	l_1	k_1
j_2	l_1	k_1
j_3	l_1	k_1
$null$	l_2	k_2

- repetition of information (e.g., the relationship l_1, k_1)
 - $(i_ID, dept_name)$
- need to use null values (e.g., to represent the tuple l_2, k_2 where there is no corresponding value for J)
 - Or do not represent this tuple at all



Testing for 3NF

- Optimization: Need to check only FDs in F , need not check all FDs in F^+ .
- Use attribute closure to check (for each dependency $\alpha \rightarrow \beta$) whether α is a superkey.
- If α is not a superkey, we have to verify if each attribute in β is contained in a candidate key of R
 - this test is rather more expensive, since it involves finding candidate keys
 - testing for 3NF has been shown to be NP-hard
 - Interestingly, **decomposition** into third normal form (described shortly) can be done in polynomial time



3NF Decomposition Algorithm

```
Let  $F_c$  be a canonical cover for  $F$ ;  
   $i := 0$ ;  
  for each functional dependency  $\alpha \rightarrow \beta$  in  $F_c$  do  
    if none of the schemas  $R_j$ ,  $1 \leq j \leq i$  contains  $\alpha \beta$   
      then begin  
         $i := i + 1$ ;  
         $R_i := \alpha \beta$   
      end  
  if none of the schemas  $R_j$ ,  $1 \leq j \leq i$  contains a candidate key for  $R$   
    then begin  
       $i := i + 1$ ;  
       $R_i :=$  any candidate key for  $R$ ;  
    end  
  /* Optionally, remove redundant relations */  
  repeat  
    if any schema  $R_j$  is contained in another schema  $R_k$   
      then /* delete  $R_j$  */  
         $R_j = R_i$ ;  
         $i = i - 1$ ;  
  return  $(R_1, R_2, \dots, R_i)$ 
```



3NF Decomposition Algorithm (Cont.)

- Above algorithm ensures:
 - each relation schema R_i is in 3NF
 - decomposition is dependency preserving and lossless-join



3NF Decomposition: An Example

- Relation schema:

cust_banker_branch = (*customer_id*, *employee_id*, *branch_name*, *type*)

- The functional dependencies for this relation schema are:

1. *customer_id*, *employee_id* \rightarrow *branch_name*, *type*
2. *employee_id* \rightarrow *branch_name*
3. *customer_id*, *branch_name* \rightarrow *employee_id*

- We first compute a canonical cover

- *branch_name* is extraneous in the r.h.s. of the 1st dependency

1. compute α^+ using only the dependencies in
$$F' = (F - \{\alpha \rightarrow \beta\}) \cup \{\alpha \rightarrow (\beta - A)\},$$
2. check that α^+ contains *A*; if it does, *A* is extraneous in β

- No other attribute is extraneous, so we get $F_C =$

customer_id, *employee_id* \rightarrow *type*
employee_id \rightarrow *branch_name*
customer_id, *branch_name* \rightarrow *employee_id*



3NF Decomposition Example (Cont.)

- The **for** loop generates following 3NF schema:

(customer_id, employee_id, type)

(employee_id, branch_name)

(customer_id, branch_name, employee_id)

- Observe that *(customer_id, employee_id, type)* contains a candidate key of the original schema, so no further relation schema needs be added
- At end of for loop, detect and delete schemas, such as (*employee_id*, *branch_name*), which are subsets of other schemas
 - result will not depend on the order in which FDs are considered
- The resultant 3NF schema is:

(customer_id, employee_id, type)

(customer_id, branch_name, employee_id)



Comparison of BCNF and 3NF

- It is always possible to decompose a relation into a set of relations that are in 3NF such that:
 - the decomposition is lossless
 - the dependencies are preserved
- It is always possible to decompose a relation into a set of relations that are in BCNF :
 - such that the decomposition is lossless
 - the dependencies may not be preserved