



# Chapter 14: Transactions

**Edited by Radhika Sukapuram**

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Conflicting Instructions

- Let  $I_i$  and  $I_j$  be two Instructions of transactions  $T_i$  and  $T_j$  respectively.
- Instructions  $I_i$  and  $I_j$  **conflict** if and only if
  - there exists some item  $Q$  accessed by both  $I_i$  and  $I_j$ ,
  - and at least one of these instructions wrote  $Q$ .
    1.  $I_i = \text{read}(Q)$ ,  $I_j = \text{read}(Q)$ .  $I_i$  and  $I_j$  don't conflict.
    2.  $I_i = \text{read}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict.
    3.  $I_i = \text{write}(Q)$ ,  $I_j = \text{read}(Q)$ . They conflict
    4.  $I_i = \text{write}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict
- Intuitively, a conflict between  $I_i$  and  $I_j$  forces a (logical) temporal order between them.
  - If  $I_i$  and  $I_j$  are consecutive in a schedule and they do not conflict,
    - ▶ their results would remain the same even if they had been interchanged in the schedule.



# Conflict Serializability

- If a schedule  $S$  can be transformed into a schedule  $S'$ 
  - by a series of swaps of non-conflicting instructions,
  - we say that  $S$  and  $S'$  are **conflict equivalent**.
- We say that a schedule  $S$  is **conflict serializable** if
  - it is conflict equivalent to a serial schedule





# Conflict Serializability (Cont.)

- Schedule 3 can be transformed into Schedule 6 –
  - a serial schedule where  $T_2$  follows  $T_1$ ,
    - ▶ by a series of swaps of non-conflicting instructions.
  - Therefore, Schedule 3 is conflict serializable.

$T_1$	$T_2$
read (A) write (A)	
	read (A) write (A)
read (B) write (B)	
	read (B) write (B)

Schedule 3

$T_1$	$T_2$
read (A) write (A) read (B) write (B)	
	read (A) write (A) read (B) write (B)

Schedule 6



# Are schedules 1 and 2 conflict equivalent?

$T_1$	$T_2$
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Schedule 1

$T_1$	$T_2$
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Schedule 2



# Conflict Serializability (Cont.)

- Example of a schedule that is not conflict serializable:

$T_3$	$T_4$
read ( $Q$ )	write ( $Q$ )
write ( $Q$ )	

- We are unable to swap instructions in the above schedule to obtain
  - either the serial schedule  $\langle T_3, T_4 \rangle$ ,
  - or the serial schedule  $\langle T_4, T_3 \rangle$ .
- How to efficiently determine conflict serializability of a schedule ?




# Precedence Graph

- Consider some schedule of a set of transactions  $T_1, T_2, \dots, T_n$
- **Precedence graph** — a directed graph where the vertices are the transactions (names).
- We draw an arc from  $T_i$  to  $T_j$  if the two transactions conflict, and  $T_i$  accesses the data item first on which the conflict arose (RW, WR, WW)
- We may label the arc by the item that was accessed.





# Precedence graph cont.






$T_1$	$T_2$
read (A) $A := A - 50$	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	$B := B + temp$ write (B) commit



# Testing for Conflict Serializability

- A schedule is conflict serializable if and only if its precedence graph is acyclic.
- Cycle-detection algorithms exist
  - (order  $(n + e)$  where  $e$  is the number of edges,  $n$  the number of vertices.)
- If precedence graph is acyclic, the **serializability order** can be obtained by a *topological sorting* of the graph.
  - For example, a serializability order for the schedule (a) would be one of either (b) or (c)





# Transaction failures

- The discussion on serializability assumes that there are no transaction failures
- Suppose there are transaction failures
  - The effect must be undone
  - A dependent transaction must also be aborted
- What is the impact of this ?



# Recoverable Schedules

- The following schedule is not recoverable if  $T_9$  commits immediately after the read(A) operation.

$T_8$	$T_9$
read (A) write (A)	
	read (A) commit
read (B)	

- If  $T_8$  should abort,
  - $T_9$  would have read (and possibly shown to the user) an inconsistent database state.
  - Hence, a database must ensure that schedules are recoverable.
- **Recoverable schedule** — if a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ ,
  - then the commit operation of  $T_i$  **must** appear before the commit operation of  $T_j$



# Cascading Rollbacks

- **Cascading rollback** – a single transaction failure leads to a series of transaction rollbacks.
- Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

$T_{10}$	$T_{11}$	$T_{12}$
read (A) read (B) write (A)	read (A) write (A)	read (A)
abort		

- If  $T_{10}$  fails,  $T_{11}$  and  $T_{12}$  must also be rolled back.
- Can lead to the undoing of a significant amount of work



# Cascadeless Schedules

- **Cascadeless schedule** — one where for each pair of transactions  $T_i$  and  $T_j$  such that
  - $T_j$  reads a data item previously written by  $T_i$ ,
    - ▶ the commit operation of  $T_i$  appears before the read operation of  $T_j$ .
- Every cascadeless schedule is also recoverable
- It is desirable to restrict the schedules to those that are cascadeless
- Example of a schedule that is NOT cascadeless

$T_{10}$	$T_{11}$	$T_{12}$
read (A) read (B) write (A)	read (A) write (A)	
abort		read (A)



# Concurrency Control

- ❑ A database must provide a mechanism that will ensure that all possible schedules are both:
  - ❑ Conflict serializable.
  - ❑ Recoverable and preferably cascadeless
- ❑ A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
- ❑ Concurrency-control schemes tradeoff between the amount of concurrency they allow and the amount of overhead that they incur
- ❑ Tests for serializability help us understand why a concurrency control protocol is correct
- ❑ **Goal** – to develop concurrency control protocols that will ensure serializability.



# Weak Levels of Consistency

- Some applications are willing to live with weak levels of consistency, allowing schedules that are not serializable
  - E.g., a read-only transaction that wants to get an approximate total balance of all accounts
  - Such transactions need not be serializable with respect to other transactions
- Tradeoff accuracy for performance





# Weak levels of consistency cont.

- Dirty read : A transaction  $T_j$  reads an uncommitted data item previously written by  $T_i$
- Dirty write : A transaction  $T_j$  writes an uncommitted data item previously written by  $T_i$



# Weak levels of consistency cont.

- Phantom read: A transaction  $T_j$  reads a data item (a phantom) that was inserted into the database by  $T_i$ 
  - $T_{30}$ : **select count(\*)**  
**from** *instructor*  
**where** *dept\_name* = 'Physics';
  - $T_{31}$  : **insert into** *instructor*  
**values** (1111, 'Feynman', 'Physics', 94000);
  - If  $T_{30}$  reads the value written by  $T_{31}$ , in a serializable schedule,  $T_{30}$  must come after  $T_{31}$ . Else  $T_{30}$  must come before  $T_{31}$
  - Conflict is on predicates, not on the data item itself
    - ▶ Information used to find tuples must also be considered for concurrency control
  - If concurrency is performed at tuple granularity, the inserted tuple may go undetected – this is called the phantom phenomenon



# Levels of Consistency in SQL-92

- **Serializable** — default [may result in non-serializable executions in some DBMSs]
- **Repeatable read** —
  - only committed records to be read, (no dirty reads)
  - repeated reads of same record must return same value. (no non-repeatable reads)
  - However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others (phantoms allowed).
    - ▶ T1 may see some records inserted by T2, but may not see others inserted by T2
- **Read committed** — only committed records can be read (no dirty reads), but successive reads of record may return different (but committed) values (non-repeatable reads and phantom reads allowed).
- **Read uncommitted** — even uncommitted records may be read.
- **All the levels above additionally disallow dirty writes**
- Syntax: **set transaction isolation level <isolation\_level>;**



# Transaction Definition in SQL

- Data manipulation languages must include a construct for specifying the set of actions that comprise a transaction.
- In SQL99, a transaction begins with **begin**
- A transaction in SQL ends by:
  - **commit work** commits current transaction and begins a new one.
  - **rollback work** causes current transaction to abort.
- In almost all database systems, by default, every SQL statement also commits implicitly if it executes successfully
  - Implicit commit can be turned off by a database directive
    - ▶ E.g. in JDBC, `connection.setAutoCommit(false);`



# Other Notions of Serializability



# View Serializability

- Let  $S$  and  $S'$  be two schedules with the same set of transactions.  $S$  and  $S'$  are **view equivalent** if the following three conditions are met, for each data item  $Q$ ,
  1. If in schedule  $S$ , transaction  $T_i$  reads the initial value of  $Q$ , then in schedule  $S'$  also transaction  $T_i$  must read the initial value of  $Q$ .
  2. If in schedule  $S$  transaction  $T_i$  executes **read**( $Q$ ), and that value was produced by transaction  $T_j$  (if any), then in schedule  $S'$  also transaction  $T_i$  must read the value of  $Q$  that was produced by the same **write**( $Q$ ) operation of transaction  $T_j$ .
  3. The transaction (if any) that performs the final **write**( $Q$ ) operation in schedule  $S$  must also perform the final **write**( $Q$ ) operation in schedule  $S'$ .
- As can be seen, view equivalence is also based purely on **reads** and **writes** alone.



# View Serializability (Cont.)

- A schedule  $S$  is **view serializable** if it is view equivalent to a serial schedule.
- Every conflict serializable schedule is also view serializable.
- Below is a schedule which is view-serializable but *not* conflict serializable.

$T_{27}$	$T_{28}$	$T_{29}$
read ( $Q$ )	write ( $Q$ )	
write ( $Q$ )		write ( $Q$ )

- What serial schedule is above equivalent to?  $\langle T_{27}, T_{28}, T_{29} \rangle$
- Every view serializable schedule that is not conflict serializable has **blind writes (writes without prior reads)**
- The problem of checking if a schedule is view serializable falls in the class of *NP*-complete problems.



# More Complex Notions of Serializability

- The schedule below produces the same outcome as the serial schedule  $\langle T_1, T_5 \rangle$ , yet is not conflict equivalent or view equivalent to it.

$T_1$	$T_5$
read (A) $A := A - 50$ write (A)	
	read (B) $B := B - 10$ write (B)
read (B) $B := B + 50$ write (B)	
	read (A) $A := A + 10$ write (A)

- If we start with  $A = 1000$  and  $B = 2000$ , the final result is 960 and 2040
- Determining such equivalence requires analysis of operations other than read and write, ie analysis of what computations are performed.
- Conclusion : There are less stringent definitions of schedule equivalence than conflict/view equivalence





# End of Chapter 14

**Edited by Radhika Sukapuram**

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use