



Chapter 15 : Concurrency Control

Edited by Radhika Sukapuram

Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Concurrency Control

- ❑ Management of concurrently executing transactions
- ❑ Ignoring failures



Lock-Based Protocols

- A lock is a mechanism to control concurrent access to a data item
- Data items can be locked in two modes :
 1. *exclusive (X) mode*. Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
 2. *shared (S) mode*. Data item can only be read. S-lock is requested using **lock-S** instruction.
- Lock requests are made to the concurrency-control manager by the programmer. Transaction can proceed only after request is granted.



Lock-Based Protocols (Cont.)

□ Lock-compatibility matrix

	S	X
S	true	false
X	false	false

- A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions
- Any number of transactions can hold shared locks on an item,
 - But if any transaction holds an exclusive on the item no other transaction may hold any lock on the item.
- If a lock cannot be granted, the requesting transaction is made to wait till all incompatible locks held by other transactions have been released. The lock is then granted.



Lock-Based Protocols (Cont.)

- Example of a transaction performing locking:

```
 $T_2$ : lock-S(A);  
      read (A);  
      unlock(A);  
      lock-S(B);  
      read (B);  
      unlock(B);  
      display(A+B)
```

- Locking as above is not sufficient to guarantee serializability — if A and B get updated in-between the read of A and B , the displayed sum would be wrong.
- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks. Locking protocols restrict the set of possible schedules.



The Two-Phase Locking Protocol

- ❑ This protocol ensures conflict-serializable schedules.
- ❑ Phase 1: Growing Phase
 - ❑ Transaction may obtain locks
 - ❑ Transaction may not release locks
- ❑ Phase 2: Shrinking Phase
 - ❑ Transaction may release locks
 - ❑ Transaction may not obtain locks
- ❑ The protocol assures serializability.
- ❑ It can be proved that the transactions can be serialized in the order of their **lock points** (i.e., the point where a transaction acquired its final lock). Proof ?



Lock Conversions

- T8: r8(a),r8(b),w8(a)
- T9: r9(a),r9(b), display9(a+b)
 - T8 will lock a in X mode. T9 cannot read a
- Two-phase locking with lock conversions:
 - First Phase:
 - can acquire a lock-S on item
 - can acquire a lock-X on item
 - can convert a lock-S to a lock-X (upgrade)
 - Second Phase:
 - can release a lock-S
 - can release a lock-X
 - can convert a lock-X to a lock-S (downgrade)
- This protocol assures serializability.
- But relies on the programmer to insert the various locking instructions.



Deadlocks

- Consider the partial schedule

T_3	T_4
lock-x (B)	
read (B)	
$B := B - 50$	
write (B)	
	lock-s (A)
	read (A)
	lock-s (B)
lock-x (A)	

- Neither T_3 nor T_4 can make progress — executing **lock-S(B)** causes T_4 to wait for T_3 to release its lock on B , while executing **lock-X(A)** causes T_3 to wait for T_4 to release its lock on A .
- Such a situation is called a **deadlock**.
 - To handle a deadlock one of T_3 or T_4 must be rolled back and its locks released.



Deadlocks (Cont.)

- Two-phase locking *does not* ensure freedom from deadlocks.
 - T_3 and T_4 are in 2PL, but deadlocked
- In addition to deadlocks, there is a possibility of **starvation**.
- **Starvation** occurs if the concurrency control manager is badly designed. For example:
 - A transaction may be waiting for an X-lock on an item, while a sequence of other transactions request and are granted an S-lock on the same item.
 - The same transaction is repeatedly rolled back due to deadlocks.
- Concurrency control manager can be designed to prevent starvation.
 - When a transaction T_i requests for a lock on data item Q
 - ▶ , a lock is granted if no other transaction is waiting for a lock on Q that made its lock request before T_i

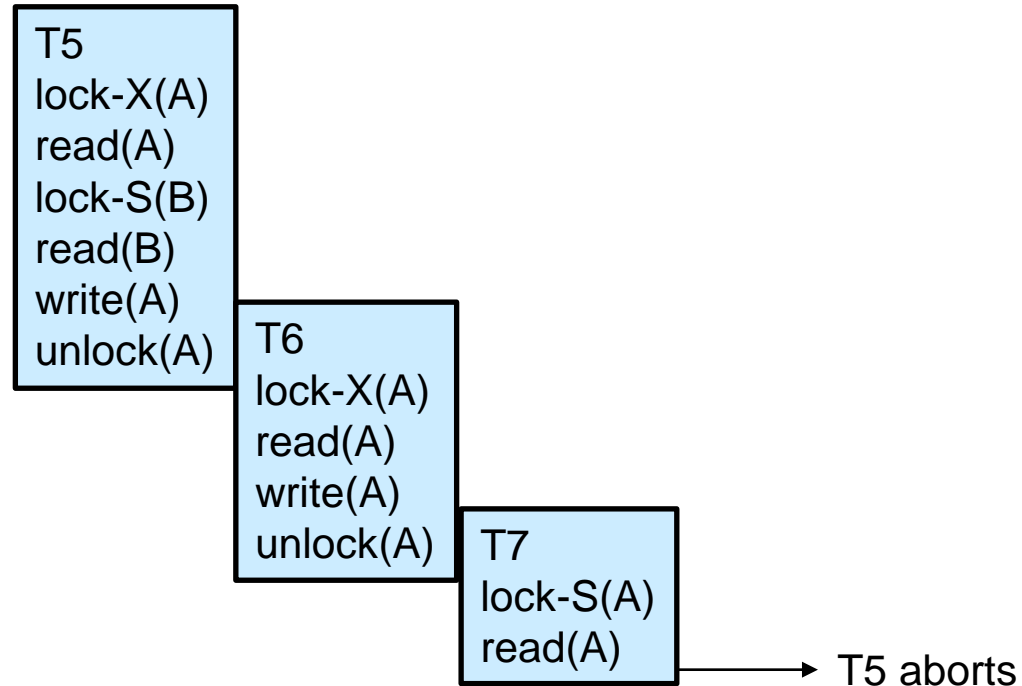


Deadlocks (Cont.)

- The potential for deadlock exists in most locking protocols. Deadlocks are a necessary evil.
- When a deadlock occurs there is a possibility of cascading roll-backs.



Strict 2PL



- ❑ Cascading roll-back is possible under two-phase locking.
 - ❑ To avoid this, follow a modified protocol called **strict two-phase locking** -- a transaction must hold all its exclusive locks till it commits/aborts.
 - ❑ Assume T1(a) releases its s-lock(a) immediately after r1(a), S2 follows strict 2PL:
 - ▶ S2: r1(a) r2(a) w2(a) commit2 commit1



Rigorous 2PL

- **Rigorous two-phase locking** is even stricter.
 - *all* locks are held till commit/abort.
 - transactions can be serialized in the order in which they commit.
 - S3: r1(a) r2(a) commit1 w2(a) commit2