



# Chapter 15 : Concurrency Control

**Edited by Radhika Sukapuram**

**Database System Concepts, 6<sup>th</sup> Ed.**

**©Silberschatz, Korth and Sudarshan**

See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Automatic Acquisition of Locks

- A transaction  $T_i$  issues the standard read/write instruction, without explicit locking calls.
- The operation **read**( $D$ ) is processed as:
  - if**  $T_i$  has a lock on  $D$
  - then**
  - read( $D$ )
  - else begin**
  - if necessary wait until no other transaction has a **lock-X** on  $D$
  - grant  $T_i$  a **lock-S** on  $D$ ;
  - read( $D$ )
  - end**



# Automatic Acquisition of Locks (Cont.)

□ **write**( $D$ ) is processed as:

if  $T_i$  has a **lock-X** on  $D$

then

write( $D$ )

else begin

if necessary wait until no other transaction has any lock on  $D$ ,

if  $T_i$  has a **lock-S** on  $D$

then

upgrade lock on  $D$  to **lock-X**

else

grant  $T_i$  a **lock-X** on  $D$

write( $D$ )

end;

□ All locks are released after commit or abort




# Implementation of Locking

- A **lock manager** can be implemented as a separate process to which transactions send lock and unlock requests
- The lock manager replies to a lock request by sending a lock grant message (or a message asking the transaction to roll back, in case of a deadlock)
- The requesting transaction waits until its request is answered
- The lock manager maintains a data-structure called a **lock table** to record granted locks and pending requests
- The lock table is usually implemented as an in-memory hash table indexed on the name of the data item being locked



# Lock Table



- Dark blue rectangles indicate granted locks; light blue indicate waiting requests
- Lock table also records the type of lock granted or requested
- New request is added to the end of the queue of requests for the data item, and granted if it is compatible with all earlier locks
- Unlock requests result in the request being deleted, and later requests are checked to see if they can now be granted
- If transaction aborts, all waiting or granted requests of the transaction are deleted
  - lock manager may keep a list of locks held by each transaction, to implement this efficiently



# Deadlock Handling

- System is deadlocked if there exists a set of waiting transactions  $\{T_0, T_1, \dots, T_n\}$  such that  $T_0$  is waiting for an item that  $T_1$  holds,  $T_1$  is waiting for an item that  $T_2$  holds, and ...,  $T_{n-1}$  is waiting for a data item that  $T_n$  holds and  $T_n$  is waiting for a data item that  $T_0$  holds.
- **Deadlock prevention** protocols ensure that the system will *never* enter into a deadlock state. A prevention strategy :
  - Require that each transaction locks all its data items before it begins execution (pre-declaration). All items are locked in one step or none are locked
    - ▶ Hard to know all data items before-hand
    - ▶ Data-item utilization may be low as they are locked for a long time



# More Deadlock Prevention Strategies

- Following schemes use transaction timestamps for the sake of deadlock prevention alone.
- $T_i$  and  $T_k$  request a data item held by  $T_j$
- **wait-die** scheme — non-preemptive
  - $i < j$ : older transaction waits for younger one to release data item
  - older means smaller timestamp
  - $k > j$ : younger transactions never wait for older ones; they are rolled back instead.
  - a transaction may die several times before acquiring the needed data item
- **wound-wait** scheme — preemptive
  - $i < j$ : older transaction *wounds* (forces rollback of) a younger transaction instead of waiting for it.
  - $k > j$ : younger transactions wait for older ones.
- Locking is used for concurrency control



# Deadlock prevention (Cont.)

- Both in *wait-die* and in *wound-wait* schemes
  - a rolled back transaction is restarted with its original timestamp. Older transactions thus have precedence over newer ones,
    - ▶ starvation is hence avoided.
- **Timeout-Based Schemes:**
  - A transaction waits for a lock only for a specified amount of time.
  - If the lock has not been granted within that time, the transaction is rolled back and restarted
    - ▶ Thus, deadlocks are not possible
    - ▶ Simple to implement, but starvation is possible.
    - ▶ Also difficult to determine a good value for the timeout interval.






# Deadlock Detection


- Deadlocks can be described as a *wait-for graph*, which consists of a pair  $G = (V, E)$ ,
  - $V$  is a set of vertices (all the transactions in the system)
  - $E$  is a set of edges; each element is an ordered pair  $T_i \rightarrow T_j$ .
- If  $T_i \rightarrow T_j$  is in  $E$ , then there is a directed edge from  $T_i$  to  $T_j$ , implying that  $T_i$  is waiting for  $T_j$  to release a data item.
- When  $T_i$  requests a data item currently being held by  $T_j$ , then the edge  $T_i \rightarrow T_j$  is inserted in the wait-for graph. This edge is removed only when  $T_j$  is no longer holding a data item needed by  $T_i$ .
- The system is in a deadlock state if and only if the wait-for graph has a cycle. Must invoke a deadlock-detection algorithm periodically to look for cycles.



# Deadlock Detection (Cont.)



Wait-for graph without a cycle



Wait-for graph with a cycle



# Deadlock Recovery

- When deadlock is detected :
  - Some transaction will have to rolled back (made a victim) to break deadlock.
    - ▶ Select that transaction as victim that will incur minimum cost.
  - Rollback -- determine how far to roll back transaction
    - ▶ **Total rollback**: Abort the transaction and then restart it.
    - ▶ More effective to roll back transaction only as far as necessary to break deadlock.
  - Starvation happens if same transaction is always chosen as victim.
    - ▶ Include the number of rollbacks in the cost factor to avoid starvation



# Order of transactions

- Determine the order between transactions at execution time by the first lock that the transactions request in incompatible mode
- Select an ordering among transactions in advance. How?



# Timestamp-Based Protocols

- Each transaction is issued a timestamp when it enters the system.
  - If an old transaction  $T_i$  has time-stamp  $TS(T_i)$ , a new transaction  $T_j$  is assigned time-stamp  $TS(T_j)$  such that  $TS(T_i) < TS(T_j)$ .
- The protocol manages concurrent execution such that the time-stamps determine the serializability order.
- To assure such behavior, the protocol maintains for each data  $Q$  two timestamp values:
  - **W-timestamp**( $Q$ ) is the largest time-stamp of any transaction that executed **write**( $Q$ ) successfully.
  - **R-timestamp**( $Q$ ) is the largest time-stamp of any transaction that executed **read**( $Q$ ) successfully.



# Timestamp-Based Protocols (Cont.)

- The timestamp ordering protocol ensures that any conflicting **read** and **write** operations are executed in timestamp order.
- Suppose a transaction  $T_i$  issues a **read**( $Q$ )
  1. If  $TS(T_i) < W\text{-timestamp}(Q)$ , then  $T_i$  needs to read a value of  $Q$  that was already overwritten.
    - Hence, the **read** operation is rejected, and  $T_i$  is rolled back.
  2. If  $TS(T_i) \geq W\text{-timestamp}(Q)$ , then the **read** operation is executed, and  $R\text{-timestamp}(Q)$  is set to  $\max(R\text{-timestamp}(Q), TS(T_i))$ .



# Timestamp-Based Protocols (Cont.)

- Suppose that transaction  $T_i$  issues **write**(Q).
  1. If  $TS(T_i) < R\text{-timestamp}(Q)$ , then the value of Q that  $T_i$  is producing was needed previously, and the system assumed that that value would never be produced.
    - Hence, the **write** operation is rejected, and  $T_i$  is rolled back.
  2. If  $TS(T_i) < W\text{-timestamp}(Q)$ , then  $T_i$  is attempting to write an obsolete value of Q.
    - Hence, this **write** operation is rejected, and  $T_i$  is rolled back.
  3. Otherwise, the **write** operation is executed, and  $W\text{-timestamp}(Q)$  is set to  $TS(T_i)$ .
- If a transaction is rolled back as a result of a read or write operation, the system assigns it a new timestamp and restarts it



# Correctness of Timestamp-Ordering Protocol

- The timestamp-ordering protocol guarantees conflict serializability since all the arcs in the precedence graph are of the form:



Thus, there will be no cycles in the precedence graph

- Timestamp protocol ensures freedom from deadlock as no transaction ever waits.
- But the schedule may not be cascade-free, and may not even be recoverable.





# Recoverability and Cascade Freedom

- Problem with timestamp-ordering protocol:
  - Suppose  $T_i$  aborts, but  $T_j$  has read a data item written by  $T_i$
  - Then  $T_j$  must abort; if  $T_j$  had been allowed to commit earlier, the schedule is not recoverable.
  - Further, any transaction that has read a data item written by  $T_j$  must abort
  - This can lead to cascading rollback --- that is, a chain of rollbacks



# One solution : Limited form of locking

- Wait for data to be committed before reading it
- A commit bit associated with each data item  $Q$ ,  $C(Q)$  is set if and only if the most recent transaction to write  $Q$  has already committed
- Suppose a transaction  $T_i$  issues a **read**( $Q$ )
  1. If  $TS(T_i) < W\text{-timestamp}(Q)$ , then  $T_i$  needs to read a value of  $Q$  that was already overwritten.
    - Hence, the **read** operation is rejected, and  $T_i$  is rolled back.
  2. If  $TS(T_i) \geq W\text{-timestamp}(Q)$ , then
    1. If  $C(Q)$  is false, the transaction waits until it becomes true or the transaction that wrote  $Q$  aborts
    2. If  $C(Q)$  is true, the **read** operation is executed, and  $R\text{-timestamp}(Q)$  is set to  $\max(R\text{-timestamp}(Q), TS(T_i))$ .
- Scheduler receives a request to commit  $T_i$ ,  $C(Q)$  is set to true.
- Scheduler receives a request to abort  $T_i$  or decides to rollback  $T_i$ 
  - Any transaction waiting for this must repeat its attempt to read  $Q$



# Thomas' Write Rule

- ❑ Modified version of the timestamp-ordering protocol in which obsolete **write** operations may be ignored under certain circumstances.
- ❑ When  $T_i$  attempts to write data item  $Q$ , if  $TS(T_i) < W\text{-timestamp}(Q)$ , then  $T_i$  is attempting to write an obsolete value of  $Q$ .
  - ❑ Rather than rolling back  $T_i$  as the timestamp ordering protocol would have done, this **write** operation can be ignored.
- ❑ Otherwise this protocol is the same as the timestamp ordering protocol.
- ❑ Thomas' Write Rule allows greater potential concurrency.
  - ❑ Allows some view-serializable schedules that are not conflict-serializable.
  - ❑  $read_{27}(Q), write_{28}(Q), write_{27}(Q)$
  - ❑ This schedule is not possible under 2PL or basic time-stamp ordering



# Limited form of locking with Thomas' Write Rule

- Rules for checking for Commit bit for read(Q) apply
- Suppose that transaction  $T_i$  issues **write**(Q).
  1. If  $TS(T_i) < R\text{-timestamp}(Q)$ , then the value of Q that  $T_i$  is producing was needed previously, and the system assumed that that value would never be produced.
    - Hence, the **write** operation is rejected, and  $T_i$  is rolled back.
  2. If  $TS(T_i) < W\text{-timestamp}(Q)$ , then  $T_i$  is attempting to write an obsolete value of Q.
    - If  $C(Q)$  is true, the write is ignored and  $T_i$  proceeds
    - If  $C(Q)$  is false,  $T_i$  waits until  $C(Q)$  becomes true or the transaction that wrote Q aborts
  3. Otherwise, the **write** operation is executed, and  $W\text{-timestamp}(Q)$  is set to  $TS(T_i)$ ,  $C(Q)$  is set to false.
- Scheduler receives a request to commit  $T_i$ ,  $C(Q)$  is set to true.
- Scheduler receives a request to abort  $T_i$  or decides to rollback  $T_i$ 
  - Any transaction waiting for this must repeat its attempt to read or **write Q**

Note: From Garcia-Molina et al



# End of Chapter 15