



Chapter 11: Indexing

Edited by Radhika Sukapuram

Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



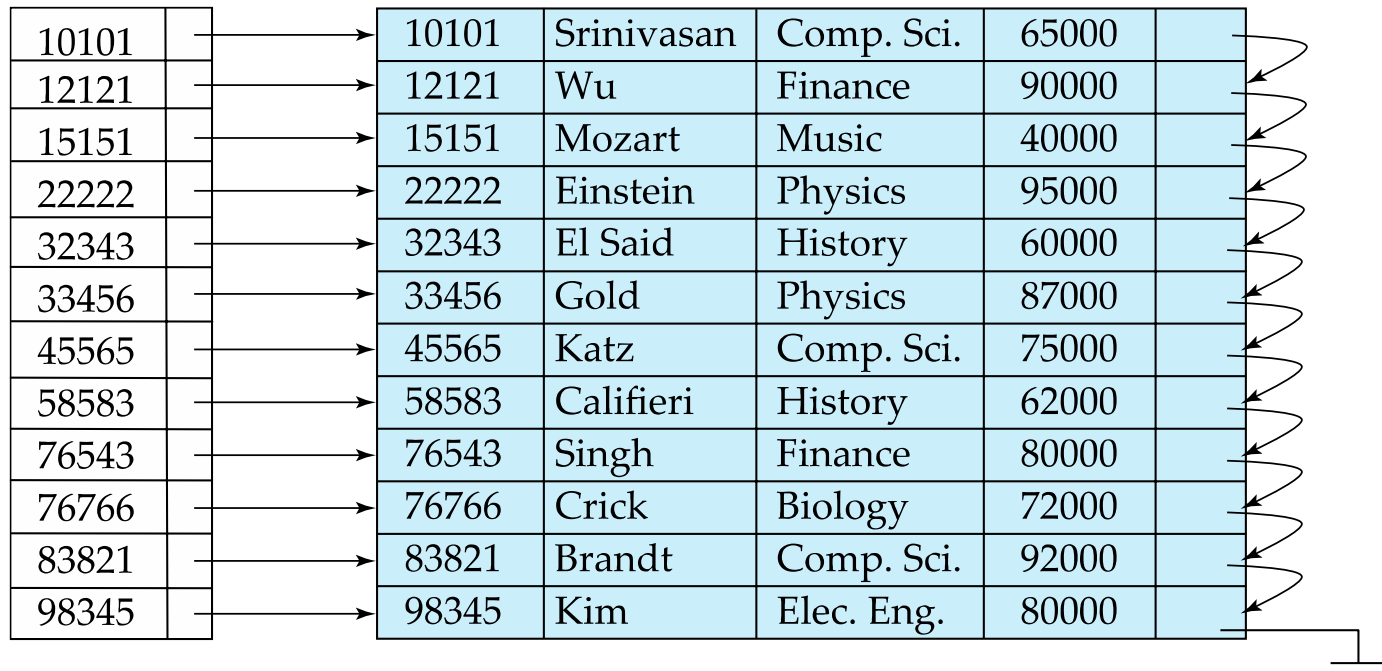
Ordered Indices

- **Ordered index:** index entries are stored based on a sorted ordering of the search key values.
 - E.g., author catalog in library.
 - **Clustering index:** in a sequentially ordered file, the index whose search key specifies the sequential order of the file.
 - ▶ Also called **primary index**
 - ▶ The search key of a clustering index is usually (but not necessarily) the primary key.
 - **Non-clustering index:** an index whose search key specifies an order different from the sequential order of the file. Also called **secondary index**.
- Two types of ordered indices: **dense** and **sparse**.



Dense Index Files

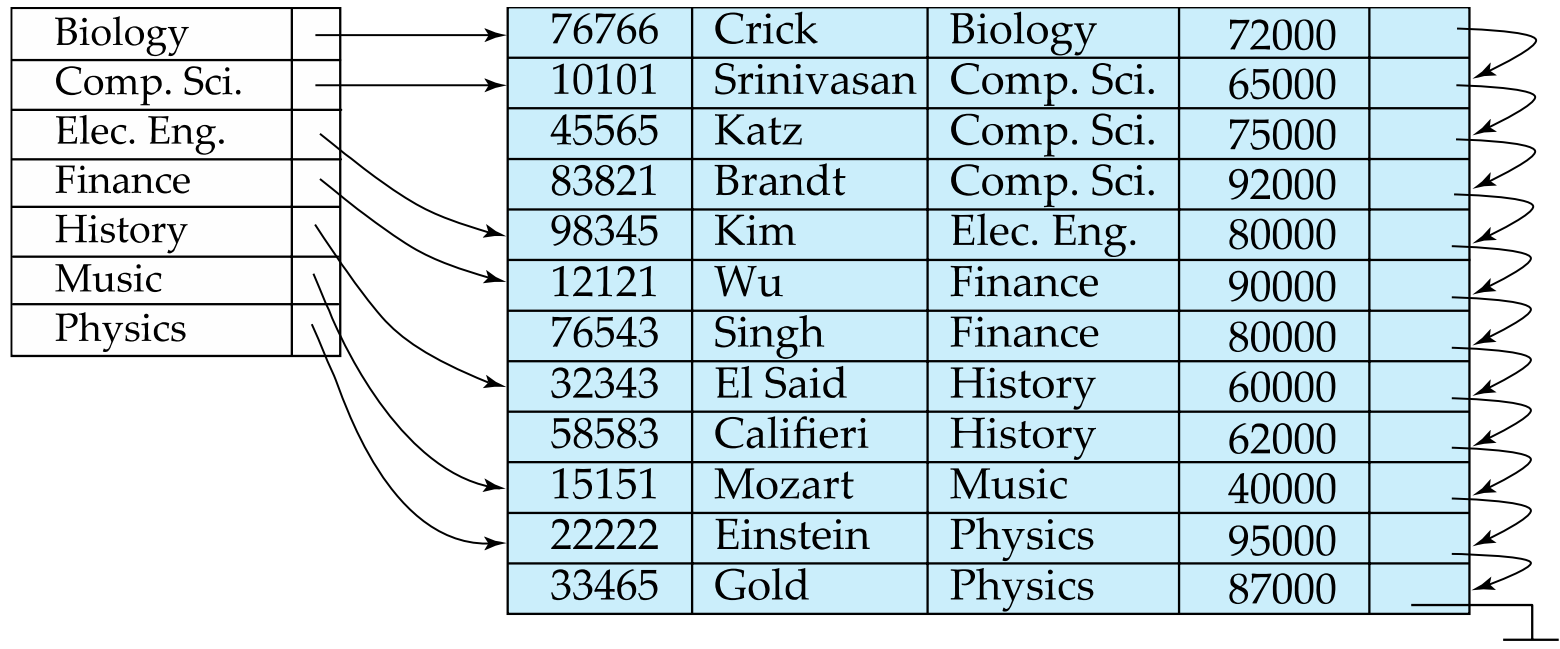
- ❑ **Dense index** — Index record appears for every search-key value in the file.
- ❑ Example: Clustering index on *ID* attribute of *instructor* relation





Dense Index Files (Cont.)

- Dense (clustering) index on *dept_name*, with *instructor* file sorted on *dept_name*

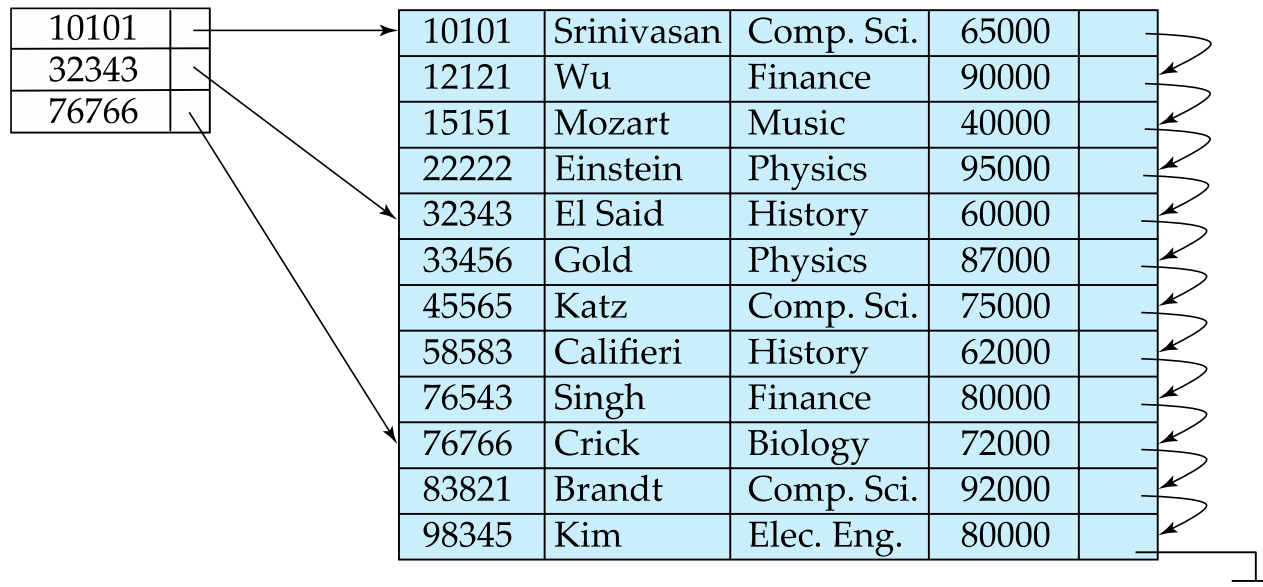


Most databases create an index on the primary key. Why ? What needs to be checked when a tuple is inserted ?



Sparse Index Files


- ❑ **Sparse Index:** contains index records for only some search-key values.
 - ❑ Applicable only when records are sequentially ordered on search-key
- ❑ To locate a record with search-key value K we:
 - ❑ Find index record with largest search-key value $< K$
 - ❑ Search file sequentially starting at the record to which the index record points





Sparse Index Files (Cont.)


- Compared to dense indices:
 - Less space and less maintenance overhead for insertions and deletions.
 - Generally slower than dense index for locating records.
- **Good tradeoff:**
 - For clustered index: sparse index with an index entry for every block in file, corresponding to least search-key value in the block.



- For un-clustered index: sparse index on top of dense index (multilevel index)



Non-clustering/secondary index



- A dense index, usually with duplicates
- Improves performance of queries that use keys other than the key of the clustering index
- A sparse secondary index is not possible. Why ?




Multilevel Index

- ❑ If a primary index does not fit in memory, access becomes expensive.
- ❑ Solution: treat primary index kept on disk as a sequential file and construct a sparse index on it.
 - ❑ Outer index – a sparse index of primary index
 - ❑ Inner index – the primary index file
- ❑ If even outer index is too large to fit in main memory, yet another level of index can be created, and so on.
- ❑ Indices at all levels must be updated on insertion or deletion from the file.



Multilevel Index Example





Deletion of Records

- ❑ **Dense indices** – If deleted record was the only record in the file with its particular search-key value,
 - ❑ the search-key is deleted from the index also
- ❑ **Sparse indices** –
 - ❑ If an entry for the search key exists in the index,
 - ▶ it is deleted by replacing the entry in the index with the next search-key value in the file (in search-key order).
 - ▶ If the next search-key value already has an index entry, the entry is deleted instead of being replaced.

10101		10101	Srinivasan	Comp. Sci.	65000	
32343		12121	Wu	Finance	90000	
76766		15151	Mozart	Music	40000	
		22222	Einstein	Physics	95000	
		32343	El Said	History	60000	
		33456	Gold	Physics	87000	
		45565	Katz	Comp. Sci.	75000	
		58583	Califieri	History	62000	
		76543	Singh	Finance	80000	
		76766	Crick	Biology	72000	
		83821	Brandt	Comp. Sci.	92000	
		98345	Kim	Elec. Eng.	80000	

Diagram illustrating the deletion of records from a file using a sparse index. The index (left table) contains three entries: 10101, 32343, and 76766. The file (right table) contains 12 records. Arrows show the mapping from index entries to file records. The record with key 32343 (El Said) is the next record in search-key order after 10101. The record with key 76766 (Crick) is the next record in search-key order after 32343. The diagram shows that the entry 32343 in the index is replaced by 76766, and the entry 76766 in the index is deleted, as it is the next search-key value already has an index entry.