

# Chapter 3: Introduction to SQL

Edited by Radhika Sukapuram. Original slides by

Database System Concepts, 6<sup>th</sup> Ed.

©Silberschatz, Korth and Sudarshan

See [www.db-book.com](http://www.db-book.com) for conditions on re-use

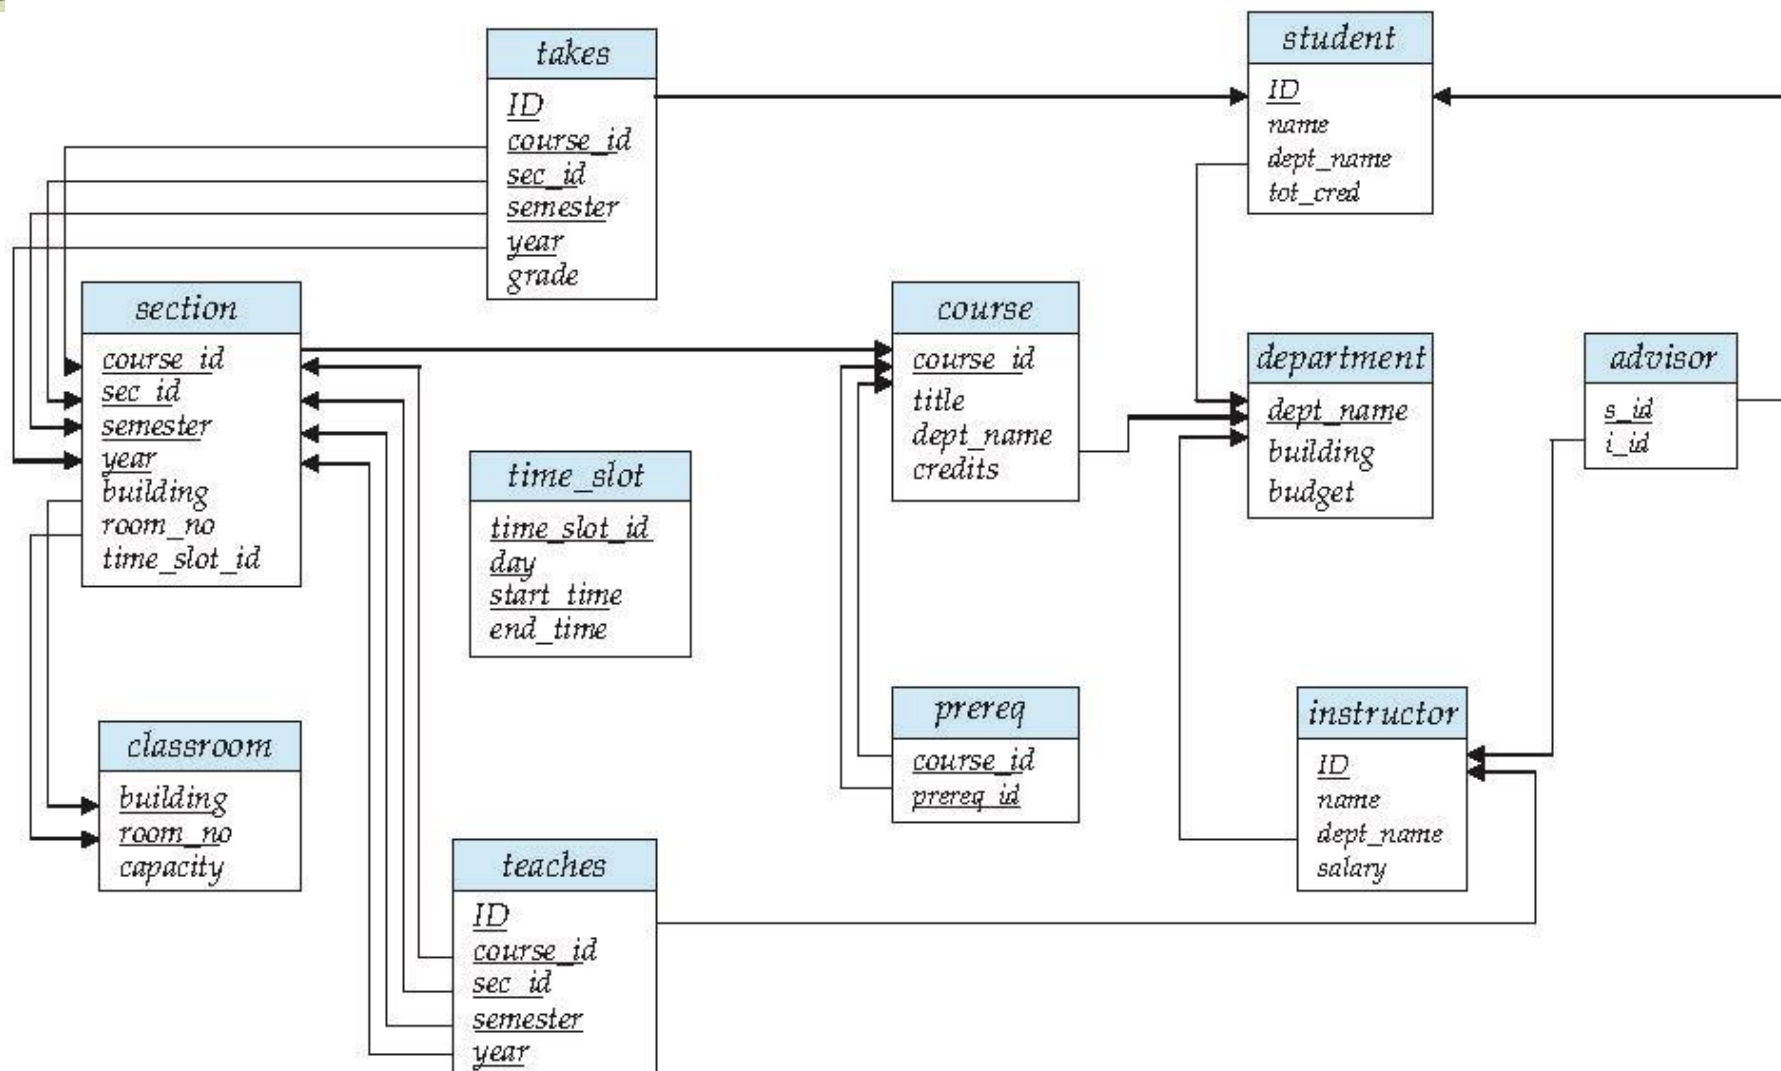


# Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
- **exists**  $r \Leftrightarrow r \neq \emptyset$
- **not exists**  $r \Leftrightarrow r = \emptyset$



# Set Operations





# Use of “exists” Clause

- Yet another way of specifying the query “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
      exists (select *
              from section as T
              where semester = 'Spring' and year = 2010
                  and S.course_id = T.course_id);
```

- Evaluate **from** clause of outer query, then predicate of outer **where** clause (partially), then select *course\_id* of that tuple
- Evaluate inner query using the *course\_id* (that is, *S.course\_id*)
- Do the same for all tuples of the outer query
  
- **Correlation name** – variable *S* in the outer query
- **Correlated subquery** – the inner query
- Scoping rules



# Use of “not exists” Clause

- Find all students who have taken all courses offered in the Biology department.



# Use of “not exists” Clause

- Find all students who have taken all courses offered in the Biology department.

```
select distinct S.ID, S.name  
from student as S  
where not exists ( (select course_id  
                    from course  
                    where dept_name = 'Biology')           12, 13, 14  
except  
                  (select T.course_id  
                   from takes as T                           12, 13, 14, 15  
                   where S.ID = T.ID));
```

- First nested query lists all courses offered in Biology
- Second nested query lists all courses a particular student took

- Note that  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$



# Test for Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- The **unique** construct evaluates to “true” if a given subquery contains no duplicates .
- Find all courses that were offered at most once in 2009

```
select T.course_id
from course as T
where unique (select R.course_id
                  from section as R
                  where T.course_id= R.course_id
                     and R.year = 2009);
```

Note: If no course is offered in 2009, the subquery returns an empty set and unique returns true

**Note: Check in latest MYSQL**



# Subqueries in the From Clause





# Subqueries in the From Clause

- ❑ SQL allows a subquery expression to be used in the **from** clause
- ❑ Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."

```
select dept_name, avg_salary
from (select dept_name, avg (salary) as avg_salary
      from instructor
      group by dept_name) as inst
where avg_salary > 42000;
```

- ❑ Note that we do not need to use the **having** clause



# With Clause

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs.
- Find all departments with the maximum budget

```
with max_budget (value) as  
    (select max(budget)  
     from department)  
select department.name  
from department, max_budget  
where department.budget = max_budget.value;
```

- Note: Check in latest MYSQL.



# Complex Queries using With Clause

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total (dept_name, value) as  
    (select dept_name, sum(salary)  
    from instructor  
    group by dept_name),  
dept_total_avg(value) as  
    (select avg(value)  
    from dept_total)  
select dept_name  
from dept_total, dept_total_avg  
where dept_total.value > dept_total_avg.value;
```

Note: Check in latest MYSQL.



# Subqueries in the Select Clause




# Scalar Subquery

- ❑ Scalar subquery is one which is used where a single value is expected
- ❑ List all departments along with the number of instructors in each department

**select** *dept\_name*,  
    (**select** **count**(\*)  
        **from** *instructor*  
        **where** *department.dept\_name* = *instructor.dept\_name*)  
    **as** *num\_instructors*  
**from** *department*;

Extracts attribute from relation



- ❑ Runtime error if subquery returns more than one result tuple
- ❑ Scalar subquery possible in **select**, **where** and **having** clauses



# Modification of the Database

- Deletion of tuples from a given relation.
- Insertion of new tuples into a given relation
- Updating of values in some tuples in a given relation



# Deletion

- Delete all instructors

**delete from** *instructor*

- Delete all instructors from the Finance department

**delete from** *instructor*  
**where** *dept\_name* = 'Finance';

- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.

**delete from** *instructor*  
**where** *dept\_name* in (**select** *dept\_name*  
                          **from** *department*  
                          **where** *building* = 'Watson');



# Deletion (Cont.)

- Delete all instructors whose salary is less than the average salary of instructors

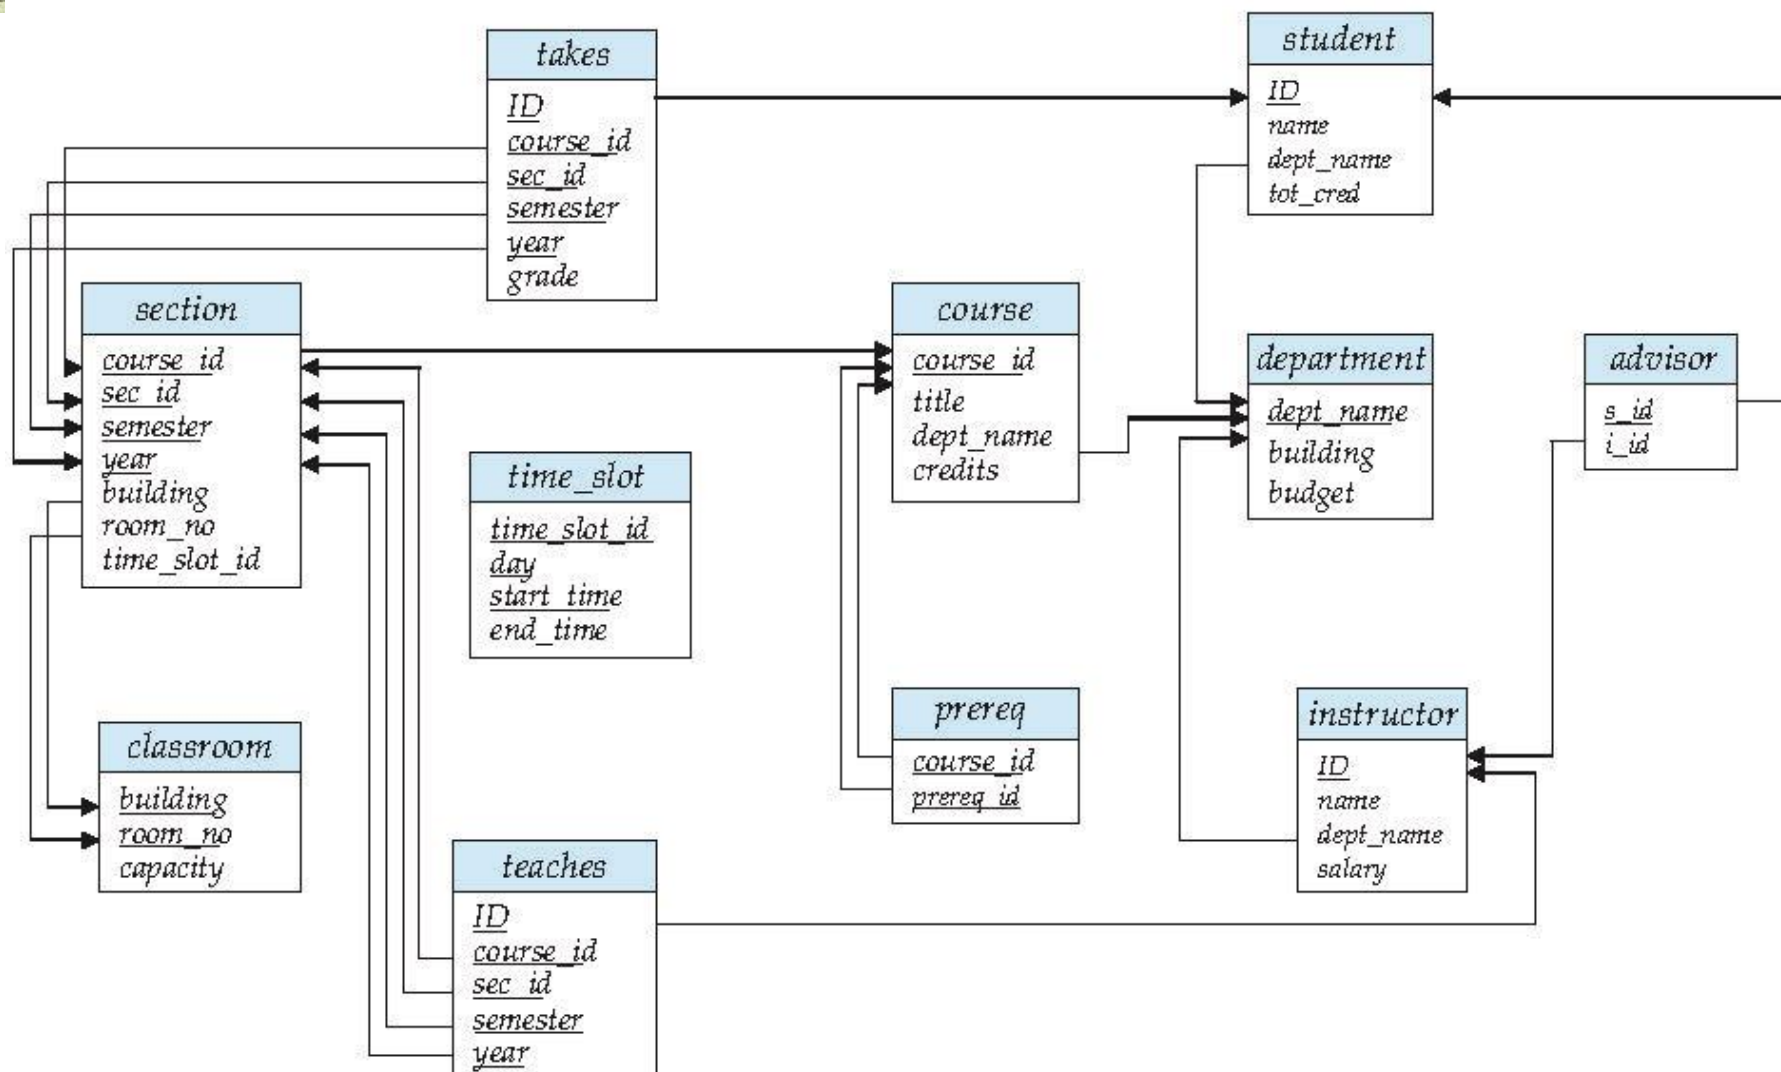
```
delete from instructor  
where salary < (select avg (salary)  
                from instructor);
```

- Problem: as we delete tuples from deposit, the average salary changes
- Solution used in SQL:
  1. First, compute **avg** (*salary*) and find all tuples to delete
  2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)





# Set Operations





# Insertion

- Add a new tuple to *course*

```
insert into course  
  values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- or equivalently

```
insert into course (course_id, title, dept_name, credits)  
  values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- Add a new tuple to *student* with *tot\_creds* set to null

```
insert into student  
  values ('3003', 'Green', 'Finance', null);
```



# Insertion (Cont.)

- Add all instructors to the *student* relation with *tot\_creds* set to 0

```
insert into student  
  select ID, name, dept_name, 0  
  from instructor
```

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation.

Otherwise queries like

```
insert into table1 select * from table1
```

would cause problems



# Updates

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
  - Write two **update** statements:  

```
update instructor
  set salary = salary * 1.03
  where salary > 100000;
update instructor
  set salary = salary * 1.05
  where salary <= 100000;
```
  - The order is important
  - Can be done better using the **case** statement (next slide)



# Case Statement for Conditional Updates

- Same query as before but with case statement

```
update instructor  
  set salary = case  
    when salary <= 100000 then salary * 1.05  
    else salary * 1.03  
  end
```



# Updates with Scalar Subqueries

- Recompute and update *tot\_creds* value for all students

**update** *student S*

**set** *tot\_cred* = (**select** **sum**(*credits*)  
**from** *takes, course*  
**where** *takes.course\_id = course.course\_id and*  
*S.ID= takes.ID.and*  
*takes.grade <> 'F' and*  
*takes.grade is not null*);

- Sets *tot\_creds* to null for students who have not taken any course

- Instead of **sum**(*credits*), use:

**case**

**when** **sum**(*credits*) **is not null** **then** **sum**(*credits*)

**else** 0

**end**

# End of Chapter 3

**Edited by Radhika Sukapuram. Original slides by  
Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use